

COMP/EECE 7/8740 Neural Networks

Topics:


- Fully connected network (FCN)
- FC layer to convolution layer
- Convolutional operations
 - Input and outputs dimensions
 - Stride size and padding
- Convolutional Neural Networks
 - Convolution layers
 - Activation layers
 - Pooling layers

Md Zahangir Alom
Department of Computer Science
University of Memphis, TN

Fully Connected Neural Network(FCNN)

FCNN consists of a series of **fully connected layers** that connect every neuron in one layer to every neuron in the other layer.

- Advantages :
 - FCNNs are “structure agnostic” (i.e. no special assumptions needed about the inputs).
- Disadvantages:
 - FCNNs show weaker performance than special-purpose networks (structure of a problem space).




Architecture : 8 (Inputs) \rightarrow 9 \rightarrow 9 \rightarrow 4 (Outputs)

Fully Connected Layer

Example: 200x200 image


40K hidden units

~2B parameters!!!



- Spatial correlation is local
- Waste of resources + we have not enough training samples in many cases


Locally Connected Layer




Example:
200x200 image
40K hidden units
Filter size: 10x10
4M parameters

Note: This parameterization is good when input image is registered (e.g., face recognition).

Locally Connected Layer



Convolutional Layer



Consider learning an image:


- Some patterns are much smaller than the whole image

Can represent a small region with fewer parameters




Same pattern appears in different places: They can be compressed!


What about training a lot of such “small” detectors
and each detector must “move around”.




Convolution




Convolution




Convolution




Convolution




Convolution




Convolution




Convolution




Convolution




Convolution




Convolution




Convolution




Convolution




Convolution




Convolution



Convolution



Convolution



Mathieu et al. “Fast training of CNNs through FFTs” ICLR 2014

Convolution

These are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

:


Each filter detects a small pattern (3 x 3).

Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Dot
product
→



6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

Convolution

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image


1	-1	-1
-1	1	-1
-1	-1	1

Filter 1




Convolution


stride=1



6 x 6 image



Filter 1



Convolution

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2


Repeat this for each filter




Two 4 x 4 images

Forming 2 x 4 x 4 matrix

Example: convolution



Color image: RGB 3 channels



Convolution: output dimension

- **Stride controls** how the filter convolves around the input volume.
- The **amount by which the filter shifts is the stride**.

Output size:


$$(N - F) / \text{stride} + 1$$

e.g. $N = 7$, $F = 3$:

$$\text{stride } 1 \Rightarrow (7 - 3)/1 + 1 = 5$$

$$\text{stride } 2 \Rightarrow (7 - 3)/2 + 1 = 3$$

$$\text{stride } 3 \Rightarrow (7 - 3)/3 + 1 = \dots$$



Example


Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5, stride 1**


Number of neurons: **5**

Output volume: **?**



Source: Jie Chen slides

Example



Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5, stride 1**

Number of neurons: **5**

$$(N - F) / \text{stride} + 1$$

Output volume: $(32 - 5) / 1 + 1 = 28$, so: **28x28x5**

How many weights for each of the 28x28x5
neurons?

Source: Jie Chen slides

Example


Examples time:

Input volume: $32 \times 32 \times 3$

Receptive fields: 5×5 , stride 2

Number of neurons: 5

Output volume: ?



Source: Jie Chen slides

Example


Examples time:

Input volume: $32 \times 32 \times 3$

Receptive fields: 5×5 , stride 2


Number of neurons: 5

Output volume: ?



Source: Jie Chen slides

Example



Examples time:

Input volume: $32 \times 32 \times 3$

Receptive fields: 5×5 , stride 2

Number of neurons: 5

$$(N - F) / \text{stride} + 1$$

Output volume: ? Cannot: $(32-5)/2 + 1 = 14.5$

Source: Jie Chen slides

Example


Examples time:

Input volume: $32 \times 32 \times 3$

Receptive fields: 5×5 , stride 3


Number of neurons: 5

Output volume: ?



Source: Jie Chen slides

Example



Examples time:

Input volume: $32 \times 32 \times 3$

Receptive fields: 5×5 , stride 3

Number of neurons: 5

$$(N - F) / \text{stride} + 1$$

Output volume: $(32 - 5) / 3 + 1 = 10$, so: $10 \times 10 \times 5$

How many weights for each of the $10 \times 10 \times 5$ neurons?

Source: Jie Chen slides

Example


Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5, stride 3**

Number of neurons: 5

Output volume: $(32 - 5) / 3 + 1 = 10$, so: **10x10x5**



Source: Jie Chen slides

In practice: Common to apply zero padding

- Padding zero in the border of images (for each channels/ feature maps)

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

neuron with receptive field 3x3, stride 1

pad with 1 pixel border => what is the output?

Output dimension with padding

- The mathematical representation with padding p as follows:

$$n_{out} = \left\lceil \frac{n_{in} + 2p - k}{s} \right\rceil + 1$$

n_{in} : number of input features

n_{out} : number of output features

k : convolution kernel size

p : convolution padding size

s : convolution stride size

In practice: Common to zero padding

0	0	0	0	0	0			
0								
0								
0								

e.g. input 7x7

neuron with receptive field 3x3, stride 1

pad with 1 pixel border => what is the output?

$$(N - F) / \text{stride} + 1 = (9-3)/1+1=7$$

7x7 => preserved size!

in general, common to see stride 1, size F, and zero-padding with $(F-1)/2$.

(Will preserve input size spatially)

Types of Convolution

“Same convolution” (preserves size)

Input [9x9]

3x3 neurons, stride 1, pad **1** => [9x9]

3x3 neurons, stride 1, pad **1** => [9x9]

- No headaches when sizing architectures
- Works well

“Valid convolution” (shrinks size)

Input [9x9]

3x3 neurons, stride 1, pad **0** => [7x7]

3x3 neurons, stride 1, pad **0** => [5x5]




- **Headaches** with sizing the full architecture
- **Works Worse!** Border information will “wash away”, since those values are only used once in the forward function


Source: Jie Chen slides

A convolutional layer

A convolutional layer has a number of filters that does convolutional operation.



Convolutional Layer



Learn multiple filters.

E.g.: 200x200 image
100 Filters
Filter size: 10x10
10K parameters

Convolution for feature extraction




Input



Feature Map


Convolution Layer

32x32x3 image -> preserve spatial structure



Convolution Layer

32x32x3 image




5x5x3 filter



Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

Convolution Layer

32x32x3 image




Filters always extend the full depth of the input volume

5x5x3 filter




Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

Convolution Layer



Convolution Layer




Convolution Layer

consider a **second, green filter**




For example, if we had $6 \times 5 \times 5$ filters, we'll get 6 separate activation maps:





We stack these up to get a “new image” of size $28 \times 28 \times 6$!

(Between, 1x1 convolution layers make perfect sense)




Convolutional Neural Networks (CNN)


- A CNN is a neural network with some convolutional layers (and some other layers).



Key operations in a CNN



Input



Feature Map

Convolution as feature extraction




Input




Feature Map

Key operations : Activation function






Rectified Linear Unit (ReLU)



$$Y = \max(0, X)$$

Source: R. Fergus, Y. LeCun

Key operation : pooling



Max

Number of feature maps don't change

Why Pooling

- Subsampling pixels will not change the object

bird



bird



We can subsample the pixels to make image smaller



fewer parameters to characterize the image

Key operations : pooling



- Max-pooling
 - partitions the input image into a set of rectangles, and for each **sub-region**, outputs the maximum value
 - Non-linear down-sampling
 - The number of output maps is the **same** as the number of input maps, but the resolution is reduced
 - Reduce the computational **complexity** for upper layers and provide a form of translation invariance
- Average pooling can also be used
- for **pool layers**, use pool size 2x2 (more = worse)

Single depth slice


1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

x
y

max pool with 2x2 filters
and stride 2



LeNet-5



- Average pooling
- Sigmoid or tanh nonlinearity
- Fully connected layers at the end
- Trained on MNIST digit dataset with 60K training examples

AlexNet model: feature visualization





Figure source: A. Karpathy

Back-prop in convolutional network


- Notes
 - https://www.cc.gatech.edu/classes/AY2018/cs7643_fall/slides/L6_cnns_backprop_notes.pdf

FC Layer Implementation




Operations in fully connected layer



CNN layer Implementation




Im2col




Im2col



CNN layer Implementation



GEMM-framework



Computation parameters

- Number of parameters in a CONV layer without bias:

$$(m * n)*k$$

- Number of parameters in a CONV layer with bias:

$$((m * n)+1)*k$$

added 1 because of the bias term for each filter.

- Here:

m: shape of width of the filter

n : shape of height of the filter

k : number of filters

Summary

- Fully connected to convolutional networks
- Convolution and sub-sampling operations
- Stride size and padding for convolution
- Convolutional Neural Networks
- LeNet architecture and AlexNet layers

- What's next?
 - How to implement better model?
 - Initialization methods
 - Activation function
 - Batch normalization
 - Regularization and
 - Optimization methods