

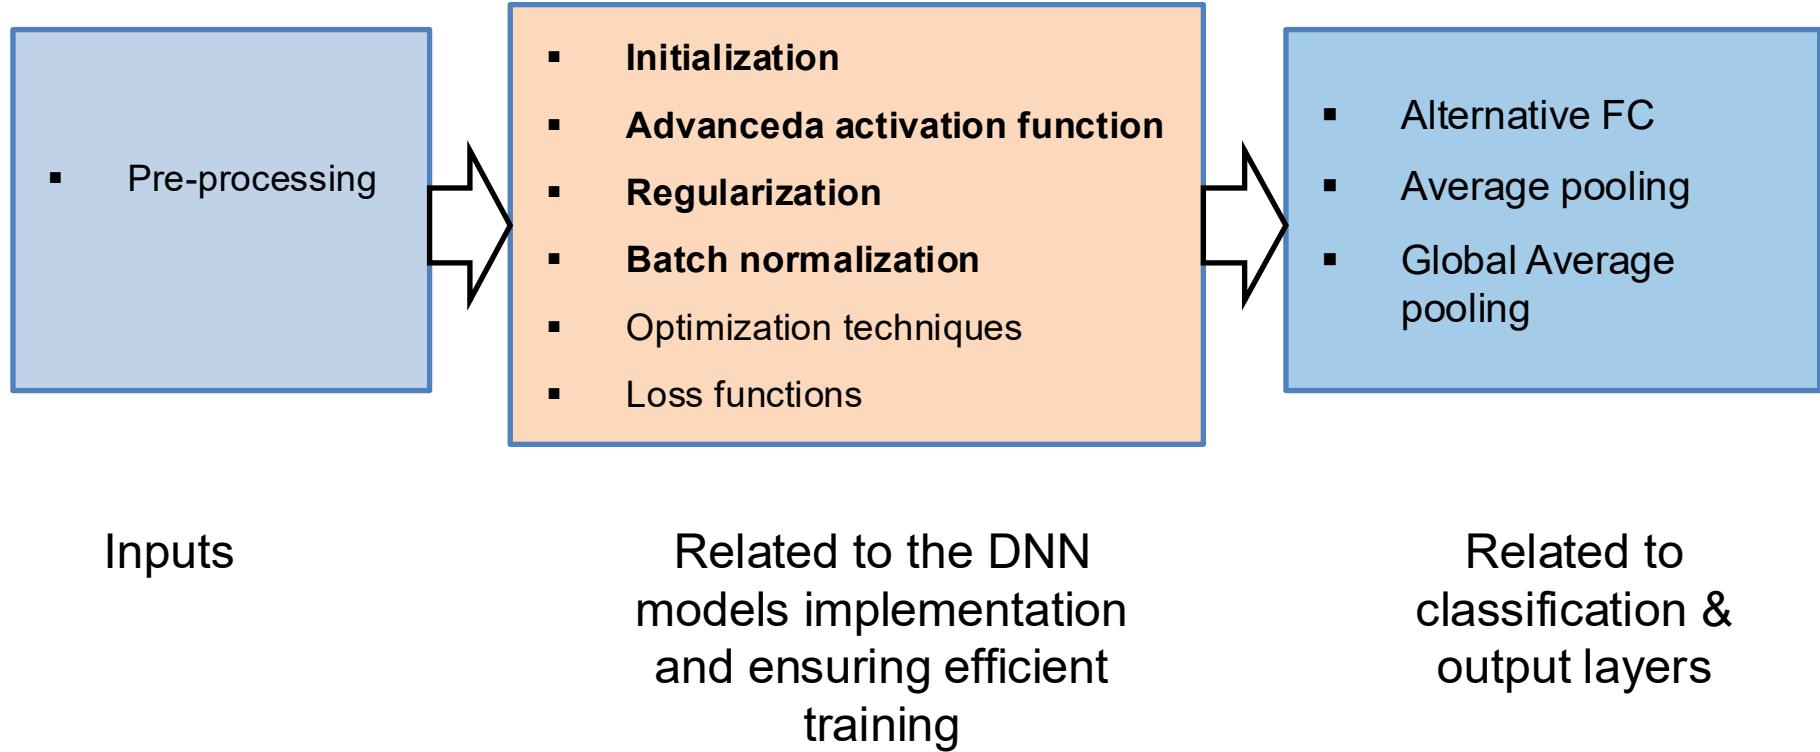
# **COMP/EECE 7/8740 Neural Networks**

## **Topics: Efficient Training Approaches**

- Input pre-processing (augmentation)
- Initialization approaches
- Advanced activation functions
- Regularization and
- Batch normalization

Md Zahangir Alom  
Department of Computer Science  
University of Memphis, TN

# Overview : Learning with DNN



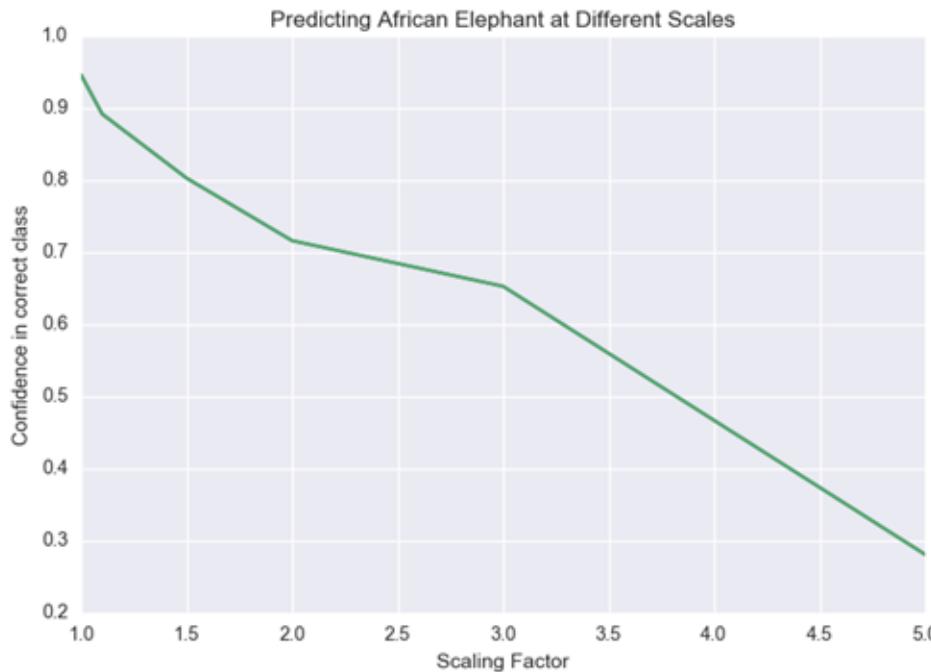
# Impact of input size

- Input this picture of an elephant to ResNet-50 at various scales
- ResNet50 was trained on 224x224 images
- Question?
  - How much bigger can I make the image before the elephant is misclassified?



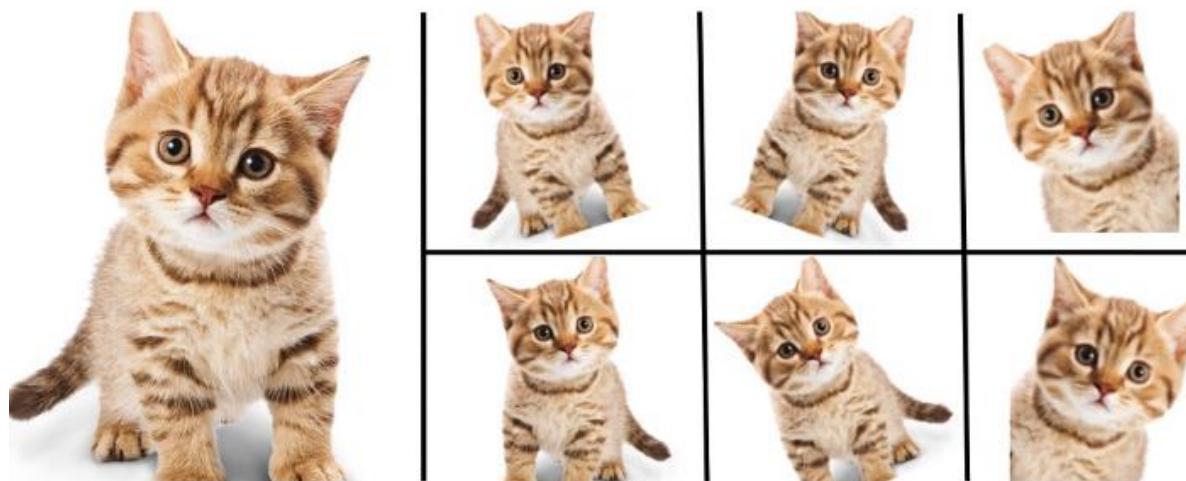
# Impact of input size

- I tried **rescales of [1.1, 1.5, 3, 5, 10]**
- Elephant was **correctly classified up till 5x scaling**
  - Input size was 1120x1120
- **Confidence of classification decays slowly**, at rescale factor of 10, ‘African Elephant’ is **no longer in the top 3**



# Data augmentation

Image **data augmentation** is a technique that can be used to **artificially expand the number** of a training dataset by **creating modified versions of images** in the dataset

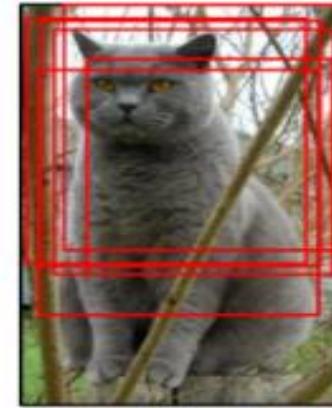


Enlarge your Dataset

# Preprocessing methods

- Sample Normalization
- Mean subtraction
- Random **cropping/ rescaling**
- Flipping sample with respective to the horizon or vertical axis
- Color jittering
- PCA/ZCA whitening
- many more. ...

- Max\_norm :  $I_{mxn} = \frac{I}{V_{max}}$
- Max\_Min\_norm :  $I_{mxn} = \frac{I - V_{min}}{V_{max} - V_{min}}$
- Rescaled :  $I_R = \frac{I}{V_{max}} * R_f$



Random cropping



Actual image



Horizontal flipping



Vertical flipping



Vertical and Horizontal flipping

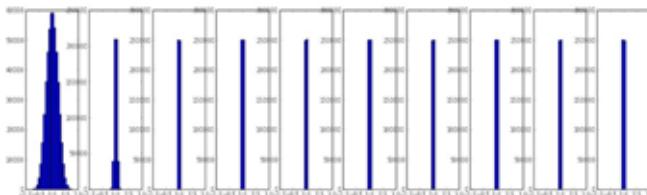
Flipping sample



Color jittering

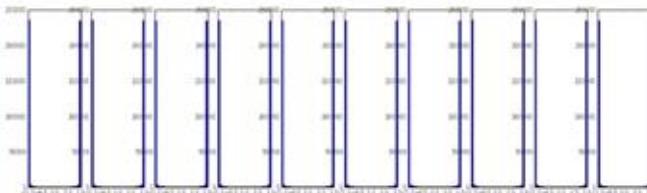
# Weight Initialization

- **Basic initialization**
- Smarter initialization schemes
- Pretraining and
- Transfer Learning



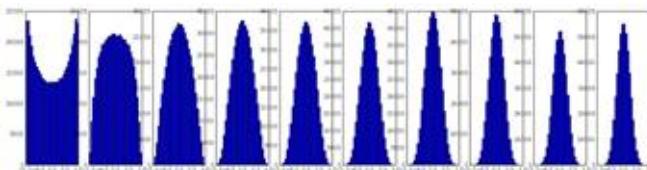
## Initialization too small:

- Activations go to zero
- gradients also zero
- No learning



## Initialization too big:

- Activations saturate (for tanh)
- Gradients zero,
- no learning



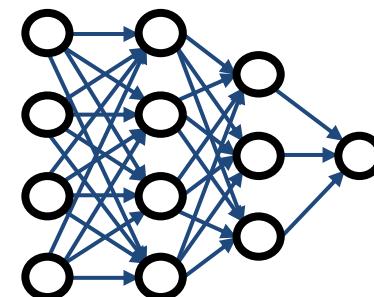
## Initialization just right:

- Nice distribution of activations at all layers,
- Learning proceeds nicely

# Baseline and smart Initialization

- **Baseline initialization approaches:** Weights cannot be initialized to same value because the **gradients will be the same**
  - Instead, **draw from some distribution**
    - Uniform from [-0.1, 0.1] is a reasonable starting spot
  - Biases may need special constant initialization
- **Smart initialization approaches:** Weights should be randomly drawn a distribution (e.g. uniform) with mean zero and standard deviation :

$$\sigma = \frac{1}{\sqrt{\text{fan-in}}}$$

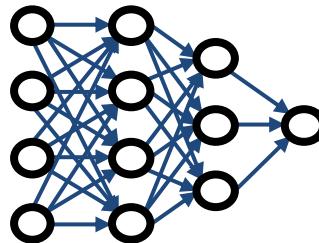


**fan\_in** the number of collections feeding into the node.

# What's Xavier initialization?

- Xavier initialization **makes sure the weights are ‘just right’**, keeping the signal in **a reasonable range of values through many layers**.
- Initialize the weights in a network will be drawn from **a distribution with zero mean and a specific variance**,

$$var(W) = \frac{1}{n_{in}}$$



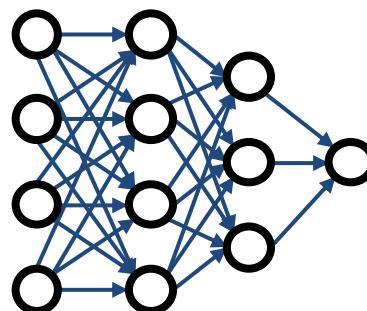
where  $W$  is the initialization distribution for the neuron in question and  **$n_{in}$  is the number of neurons feeding into it**.

The distribution used is typically Gaussian or uniform.

# Glorot & Bengio's initialization scheme

- Proper initialization of deep networks is important because of the **multiplicative effect through layers**
- This technique satisfy objectives of maintaining **activation variances and back-propagated gradients variance** as one moves up or down the network.
- We call it the **normalized initialization**:

$$W \sim U \left[ -\frac{\sqrt{6}}{\sqrt{n_j + n_{j+1}}}, \frac{\sqrt{6}}{\sqrt{n_j + n_{j+1}}} \right]$$



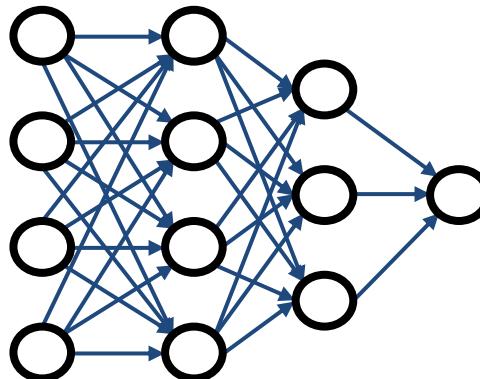
# Glorot & Bengio's initialization scheme

- For **Tanh** units: sample a Uniform  $(-r, r)$  with

$$r = \sqrt{\frac{6}{N_{in} + N_{out}}}$$

- For **sigmoid** units: sample a Uniform  $(-r, r)$  with

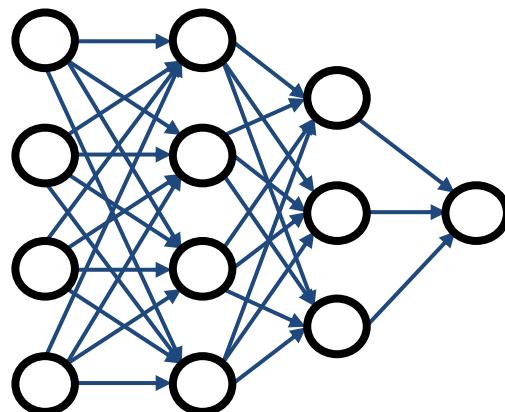
$$r = 4\sqrt{\frac{6}{N_{in} + N_{out}}}$$



# He initialization (best one)

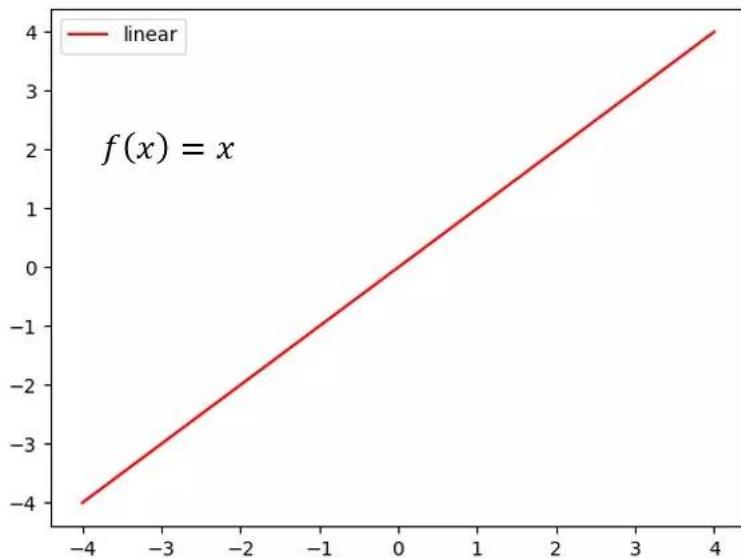
$$w_l \sim \mathcal{N}\left(0, \frac{2}{n_l}\right)$$

Number of inputs to neuron

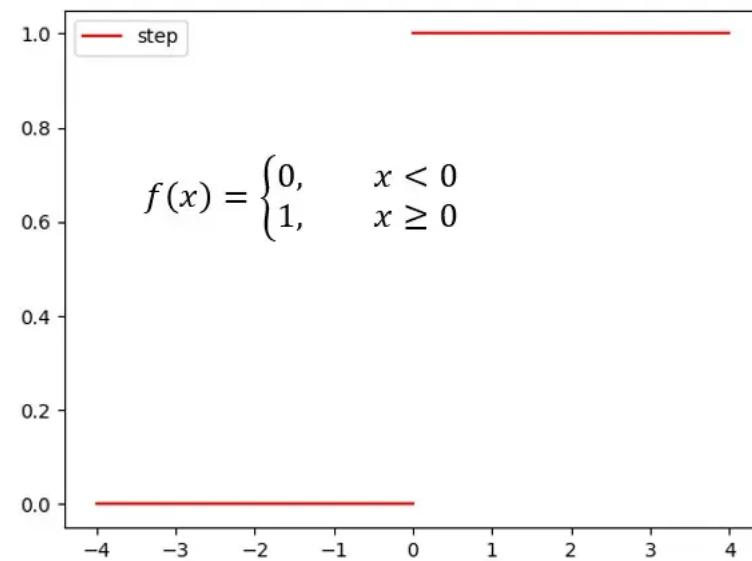


# Activation Functions

$$\vec{y} = \text{NeuralNetwork}(\vec{x}) = \sigma(\sigma(\sigma(\sigma(\vec{x} \cdot W_1) \cdot W_2) \cdot W_3) \cdot W_4)$$



Linear activation

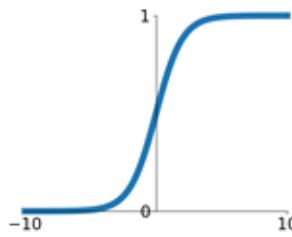


Step function

# Activation Functions

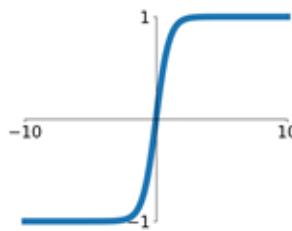
## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



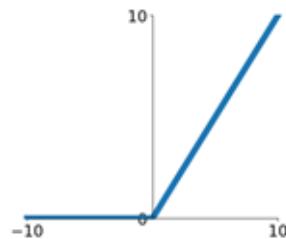
## tanh

$$\tanh(x)$$



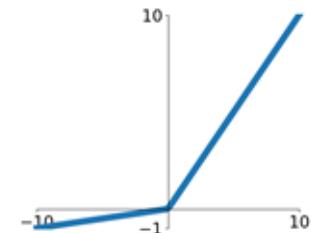
## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$

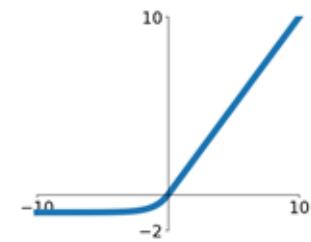


## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

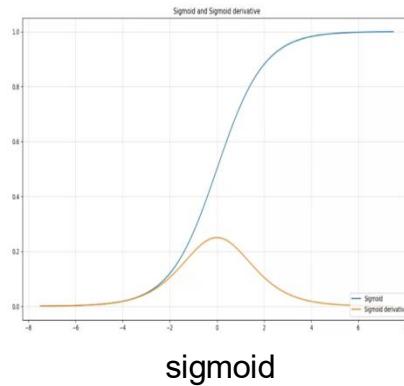
## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$

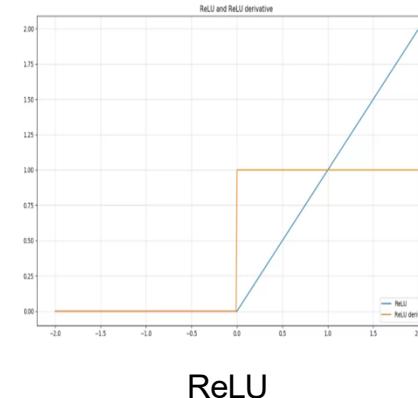


# Derivative of activation functions

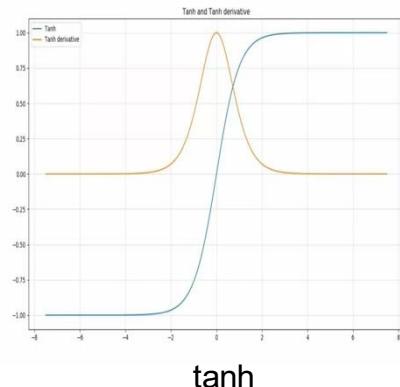
- The derivative of the **ReLU function** is one at every point above zero, creating a more stable network.



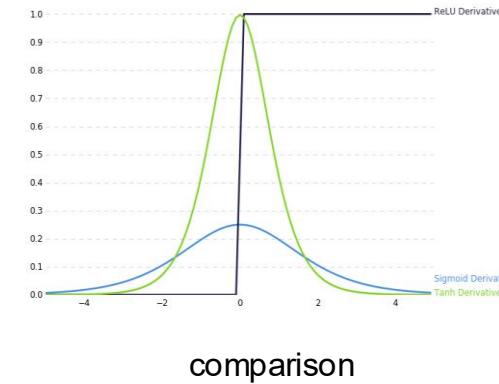
sigmoid



ReLU



tanh



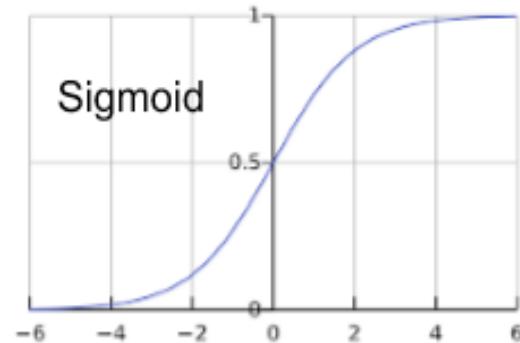
comparison

**Issues :** exploding and vanishing Gradients for Tanh and sigmoid activations !

# Problem with saturation

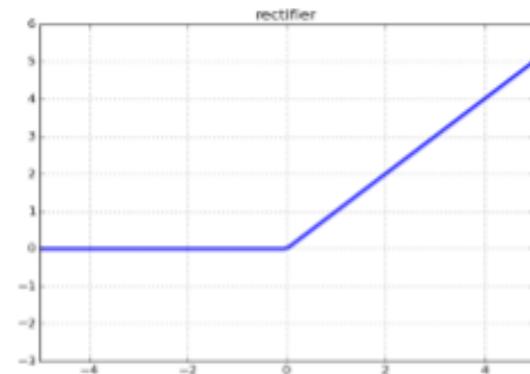
## Issue:

- A null gradient results in no learning, which happens if:
  1. Sigmoid saturates or
  2. ReLU saturates



## Solution:

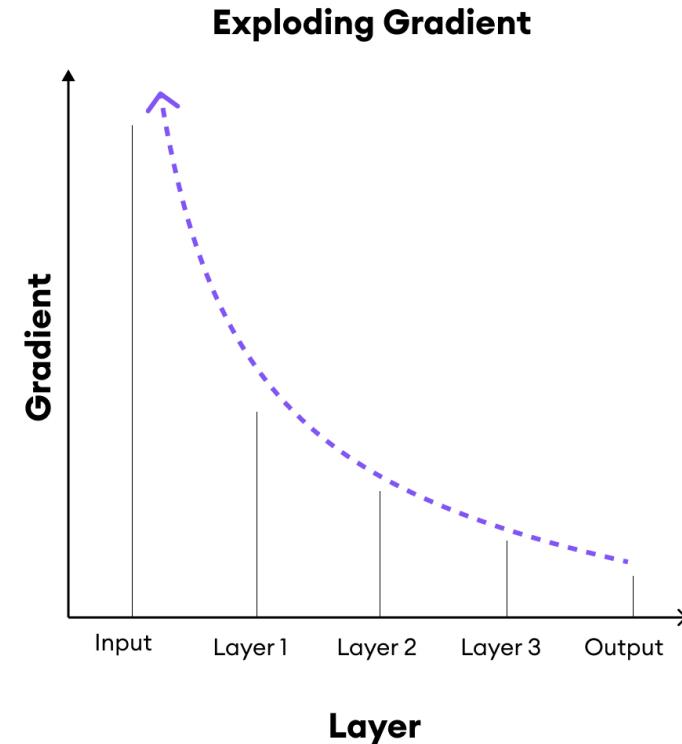
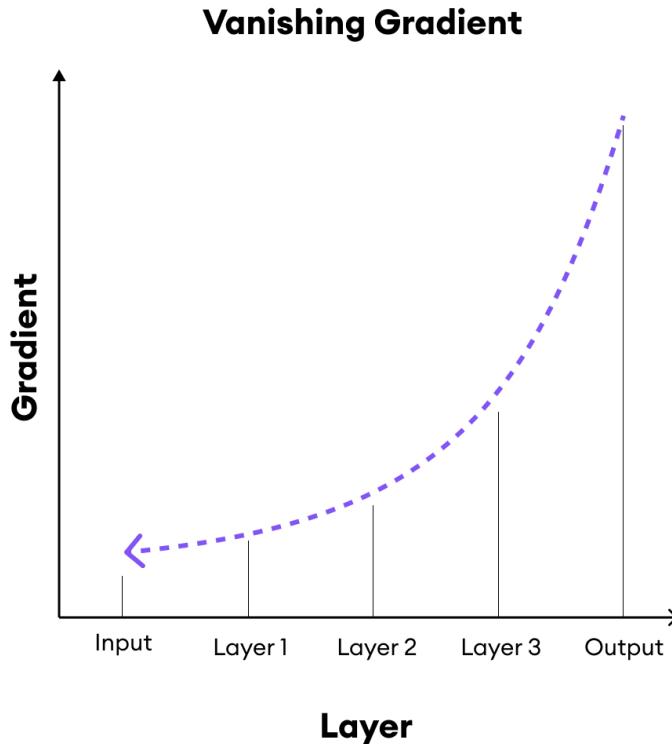
- Initialize the weights so that the average value is zero, i.e., **work in the interesting zone of activation function**
- Normalize data (zero mean)



$$RELU(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

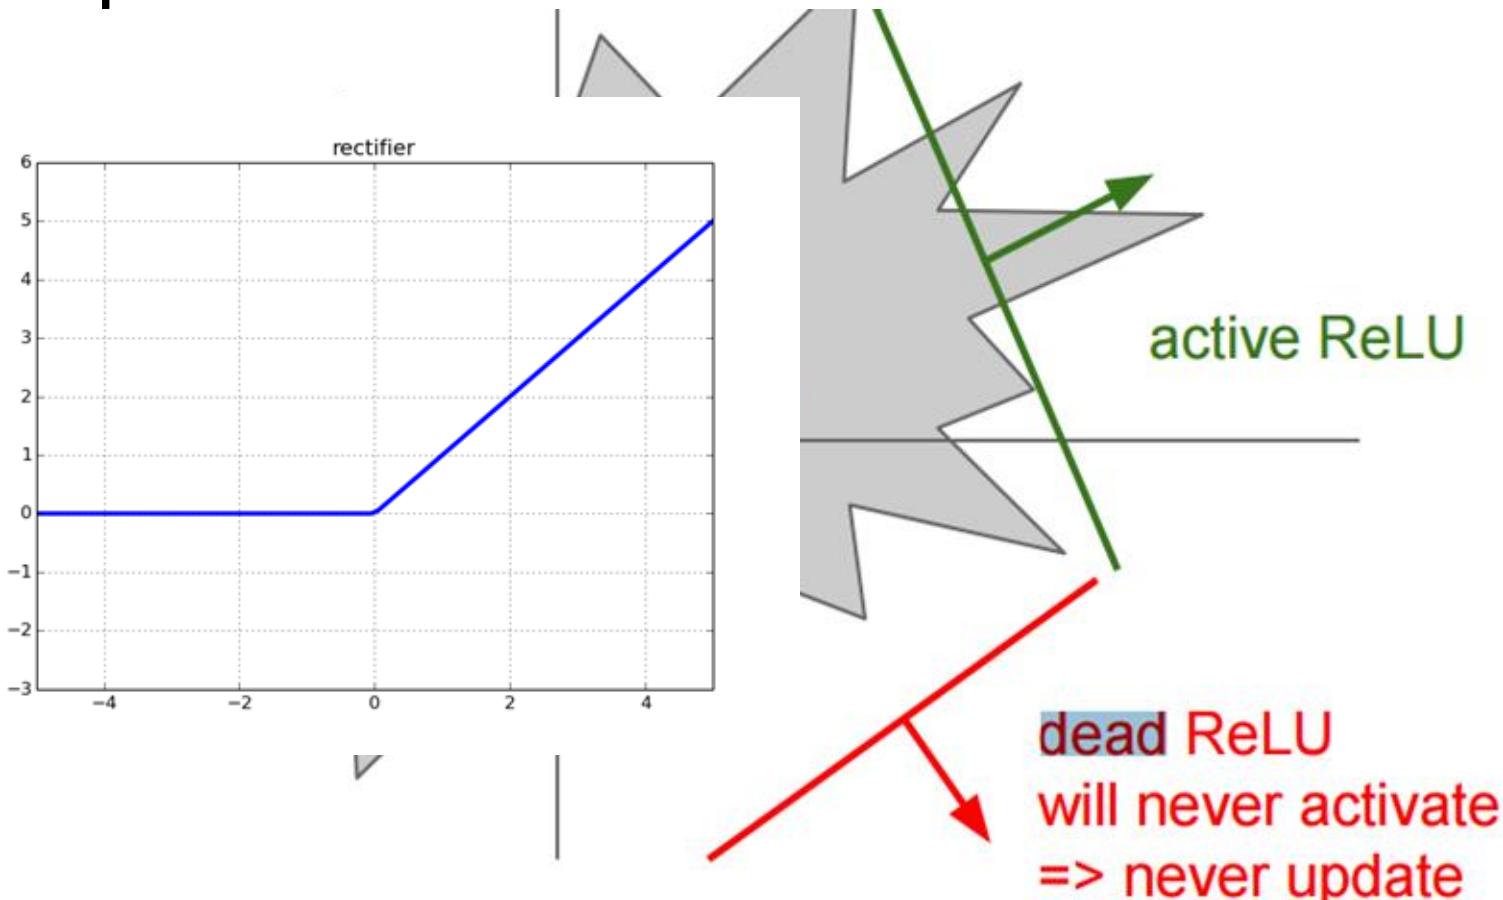
# Gradient exploding and saturation

- Gradient status within the networks (respect to the input and output layers)



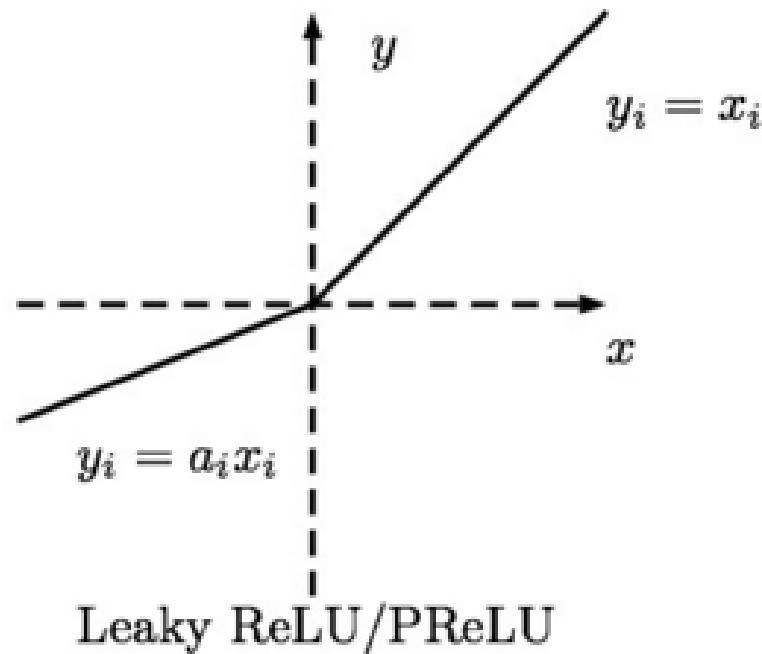
# Dying ReLU Problem

- **ReLU dies**, if input to ReLU is negative for the dataset
- Brief burst of **research into addressing dying ReLUs**
- General idea is to have **non-zero gradients even for negative inputs**



# Leaky ReLU & Parameterized ReLU

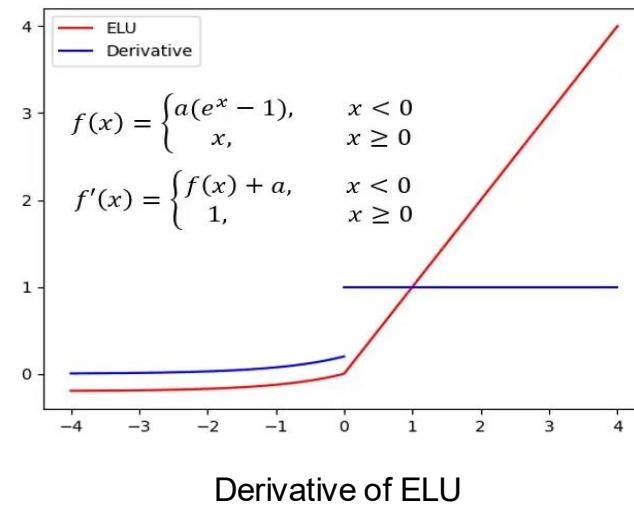
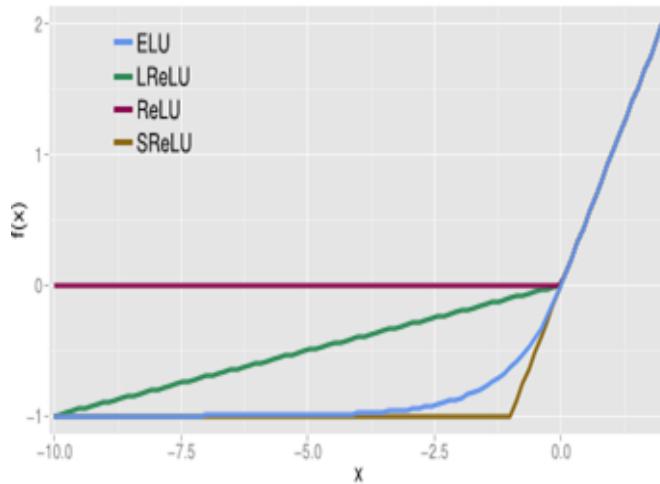
- In Leaky ReLU,  $a$  is a hyperparameter.
- In Parameterized ReLU(PReLU),  $a$  is learned.



# ELU and other activation functions

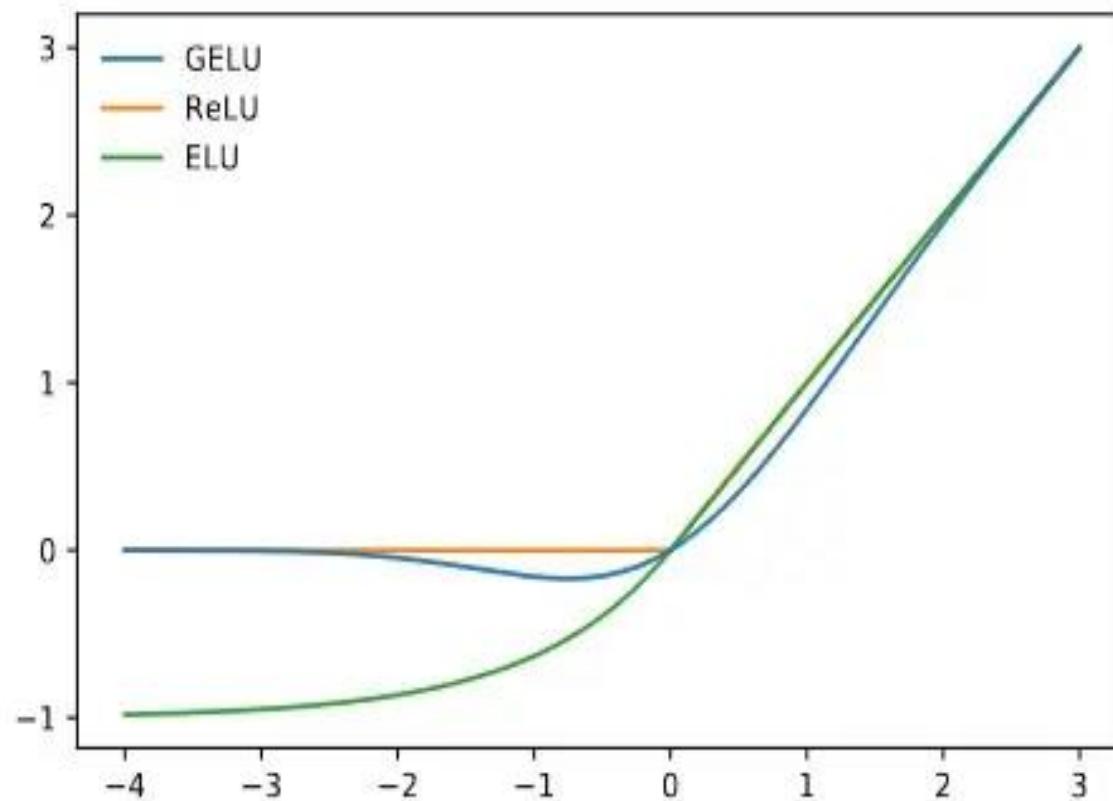
- The papers introducing each **alternative activations claim they work well**
- All the architectures, we are about to discuss used **SReLUs**

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha(\exp(x) - 1) & \text{if } x \leq 0 \end{cases}, \quad f'(x) = \begin{cases} 1 & \text{if } x > 0 \\ f(x) + \alpha & \text{if } x \leq 0 \end{cases}$$



# Gaussian ELU (GELU)

$$GELU(x) = 0.5x(1 + \tanh(\sqrt{2/\pi}(x + 0.044715x^3)))$$



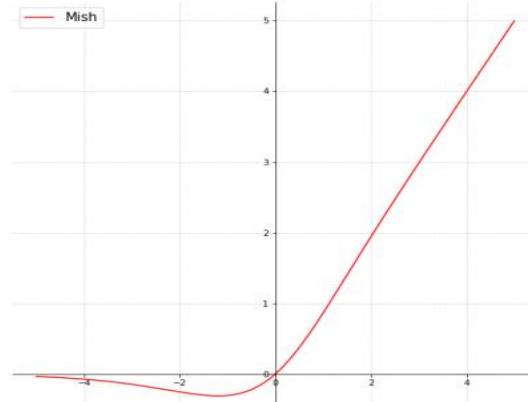
Hendrycks, Dan, and Kevin Gimpel. "Gaussian error linear units (gelus)." *arXiv preprint arXiv:1606.08415* (2016).

# Recent trend: Hybrid Activation functions

- Mish Activation function

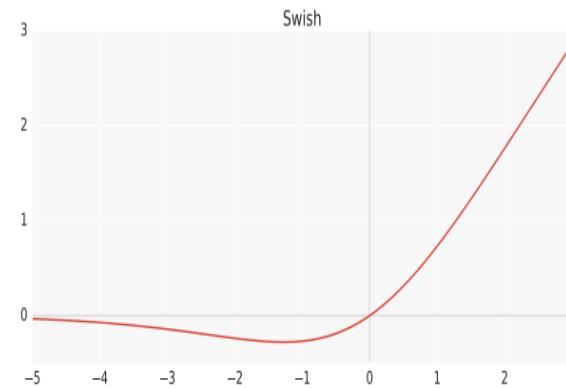
$$f(x) = x \cdot \tanh(\varsigma(x))$$

$$\varsigma(x) = \ln(1 + e^x)$$



- Swish Activation function

$$f(x) = x \cdot \text{sigmoid}(x).$$

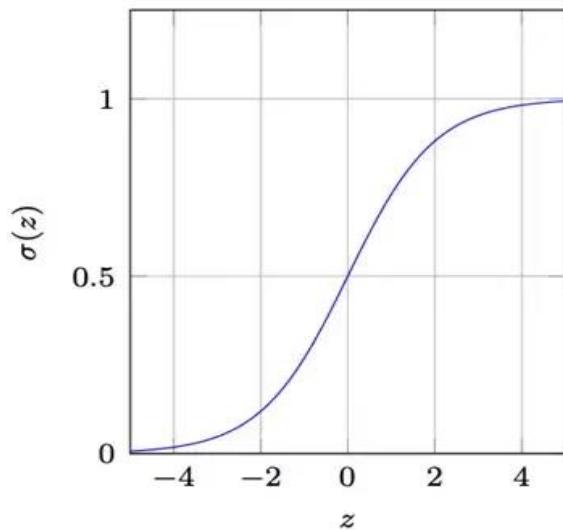


Misra, Diganta. "Mish: A Self Regularized Non-Monotonic Neural Activation Function." *arXiv preprint arXiv:1908.08681* (2019).

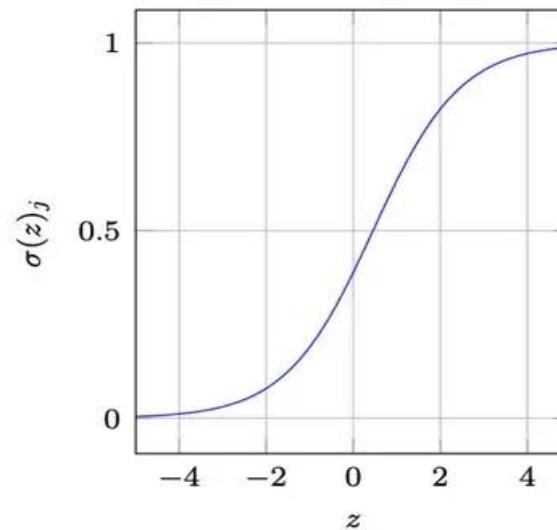
Ramachandran, Prajit, Barret Zoph, and Quoc V. Le. "Swish: a self-gated activation function." *arXiv preprint arXiv:1710.05941* 7 (2017).

# Activation function in the output layer

- We usually consider the following two activation function for the output layer:
  - Sigmoid : two (binary) classes
  - Softmax : more than two classes



(a) Sigmoid activation function.



(b) Softmax activation function.

# How to choose an activation function?

- **Million-dollar question in deep learning:** how to actually choose the right activation function when training a neural network from scratch?
- Different activation functions have different advantages and disadvantages and depending on the type of the **neural network the outcome may be different.**
- Starting point can be to choose one of the ReLU-based activation functions (including ReLU itself)[1]
  - Increase the convergence speed.
- For **classification layer:**
  - binary classification then the sigmoid activation function is a good choice
  - For the multi-class classification softmax function is better as it will output probability representation for each class.

[1] Lederer, Johannes. "Activation functions in artificial neural networks: A systematic overview." *arXiv preprint arXiv:2101.09957* (2021).

# What “same activation and gradient updating” really means

- Usually mean **consistency and compatibility between:**
  - the **activation function** used in forward propagation, and
  - how **gradients behave and are updated** during backpropagation.
- Why?:
  - **The activation function controls the gradient flow.**
  - **If gradients behave badly**, training becomes slow, unstable, or impossible.

## What can go wrong?

If derivatives are **too small** → **vanishing gradients**

If derivatives are **too large** → **exploding gradients**

# Where it helps and how?

- Avoiding biased or dead update:
  - If activation and gradient update are mismatched:
    - weights in early layers update **too slowly**
    - some neurons **never activate again** (dead neurons)
    - learning focuses only on top layers
  - This leads to:
    - slow convergence
    - poor generalization
    - wasted model capacity

# Initialization & activation work together

- **Same-scale activations → same-scale gradients**
  - Efficient learning requires:
    - activations have **zero mean**
    - controlled variance
    - gradients have similar scale across layers
- Why these pairings matter:

Activation	Initialization	Reason
Sigmoid / Tanh	Xavier	Keeps variance stable
ReLU / Leaky ReLU	He	Compensates for half-zero outputs

- If you mismatch them:
  - variance explodes or collapses
  - gradients become unstable

Consistent activation behavior → consistent gradient scale → **one learning rate works well.**

# Information preservation perspective (important intuition)

- Think of training as **information flowing forward and backward**.
  - Forward pass: data → features
  - Backward pass: error → credit assignment
- Good activation functions:
  - preserve information forward
  - preserve learning signal backward
- Bad ones:
  - squash information
  - block gradients
- Efficient training requires **symmetry in information flow**.

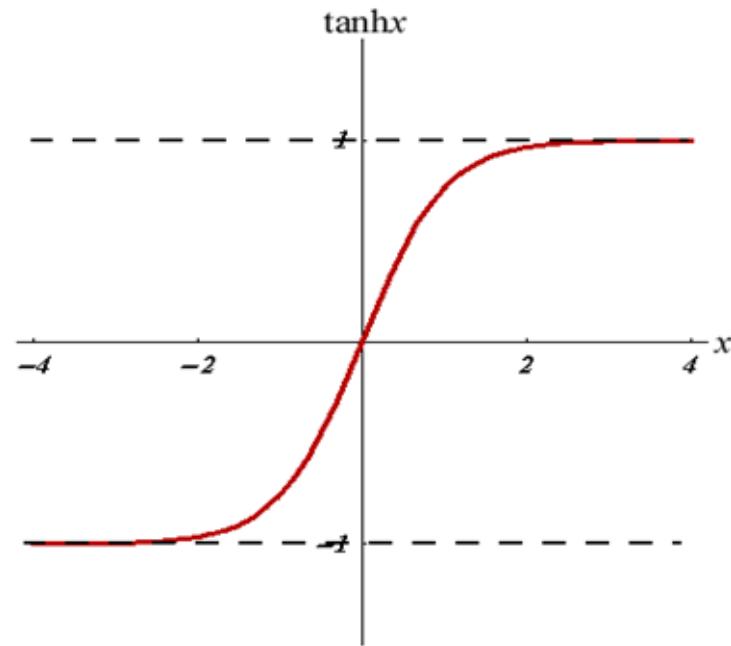
[1] Lederer, Johannes. "Activation functions in artificial neural networks: A systematic overview." *arXiv preprint arXiv:2101.09957* (2021).

# Regularization outline

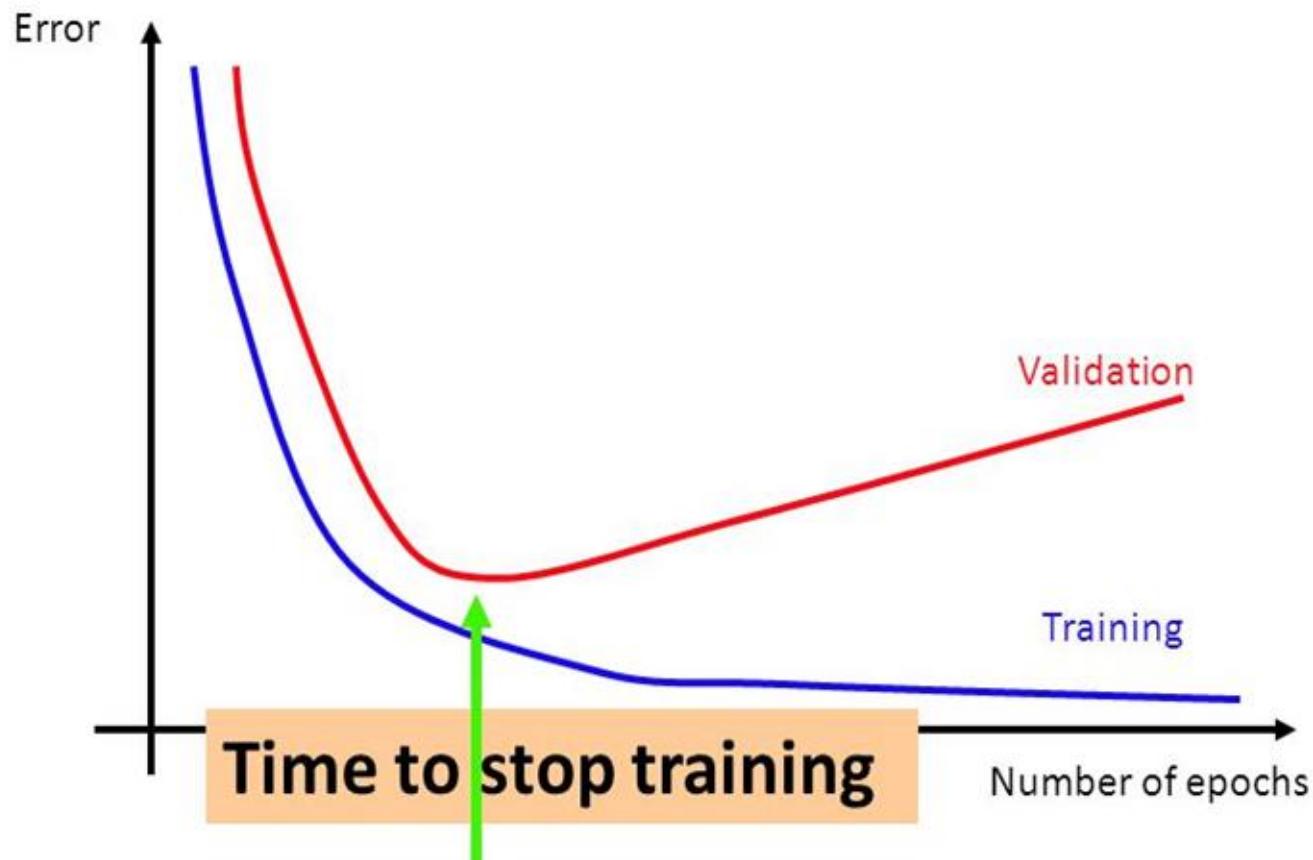
- **L1 / L2 regularization**
- Early stopping
- Auxiliary classifiers
- Penalizing confident of output distributions and
- Dropout
- Batch normalization

# L1 / L2 regularization

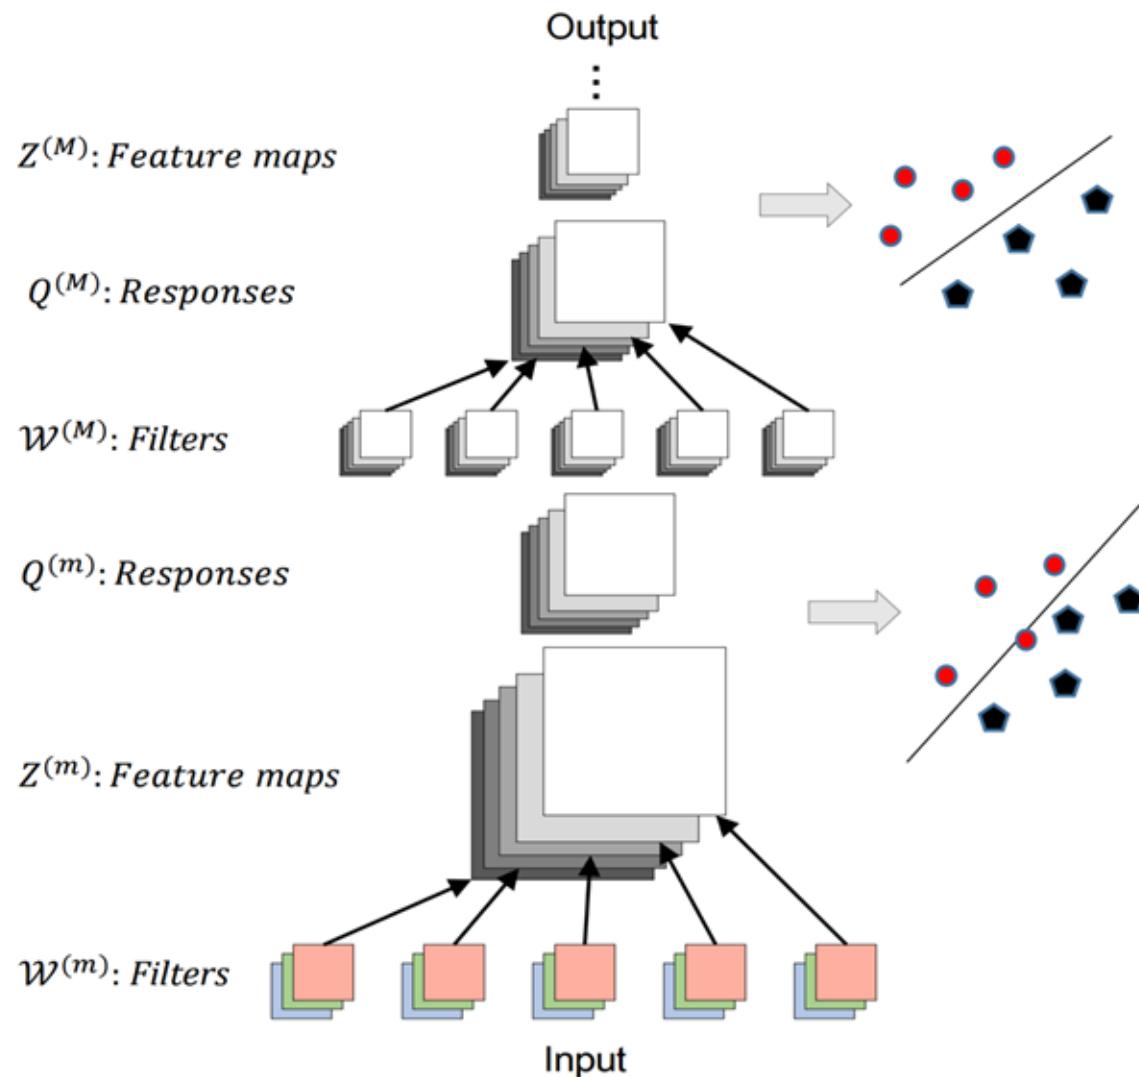
- L1 encourages sparsity
- L2 discourages large weights
  - Gaussian prior on weight



# Early Stopping



# Auxiliary Classifiers



# Penalizing outputs distributions

- Do not allow the model to be overconfident
- Prefer smoother output distribution
- Invariant to model parameterization
  1. Train towards smoother distribution
  2. Penalize entropy

True Label

$$\boxed{y_i} = [0, 0, 1, 0]$$

Baseline

$$\boxed{u} = [0.25, 0.25, 0.25, 0.25]$$

$$\hat{y}_i = (1 - \epsilon) \boxed{y_i} + \epsilon \boxed{u}$$

$$\epsilon = 0.04$$

Mixed Target

$$\hat{y}_i = [0.01, 0.01, 0.97, 0.01]$$

# Penalizing confident distributions

$$y_i = [0, 0, 1, 0]$$

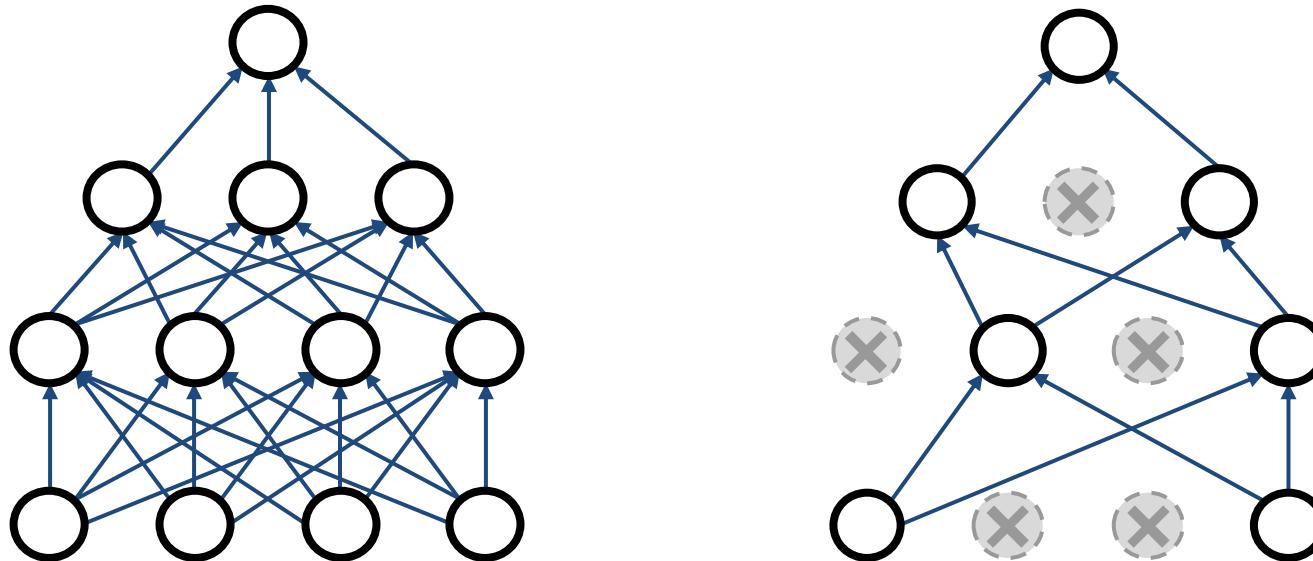
$$u = [0.25, 0.25, 0.25, 0.25]$$

$$\hat{y}_i = (1 - \epsilon)y_i + \epsilon u \quad \text{If } \epsilon = 0.04$$

$$\hat{y}_i = [0.01, 0.01, 0.97, 0.01]$$

# Regularization: Dropout

- In each forward pass, **randomly set some neurons to zero**
- Probability of dropping is a **hyperparameter**; 0.5 is common



Srivastava et al, "Dropout: A simple way to prevent neural networks from overfitting", JMLR 2014

Slide Credit: Fei-Fei Li, Justin Johnson, Serena Yeung, CS 231n

# Regularization: Dropout

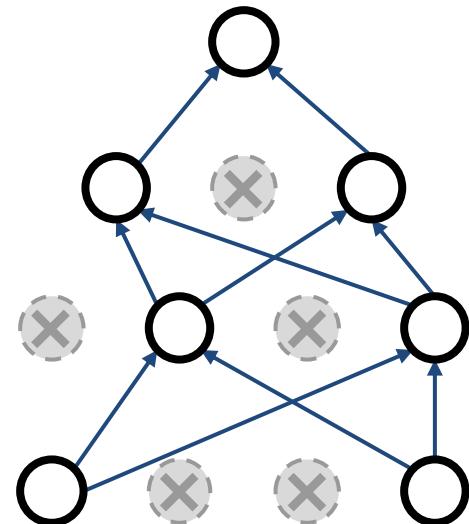
```
p = 0.5 # probability of keeping a unit active. higher = less dropout

def train_step(X):
    """ X contains the data """

    # forward pass for example 3-layer neural network
    H1 = np.maximum(0, np.dot(W1, X) + b1)
    U1 = np.random.rand(*H1.shape) < p # first dropout mask
    H1 *= U1 # drop!
    H2 = np.maximum(0, np.dot(W2, H1) + b2)
    U2 = np.random.rand(*H2.shape) < p # second dropout mask
    H2 *= U2 # drop!
    out = np.dot(W3, H2) + b3

    # backward pass: compute gradients... (not shown)
    # perform parameter update... (not shown)
```

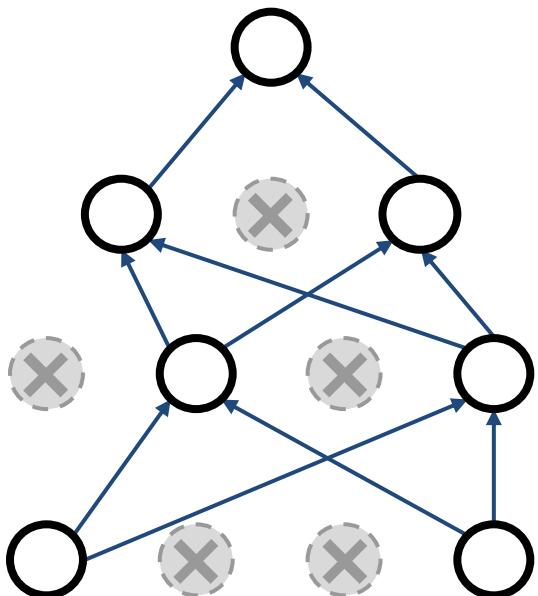
Example forward pass  
with a 3-layer network  
using dropout



# Regularization: Dropout

How can this possibly be a good idea?

Forces the network to have a redundant representation;



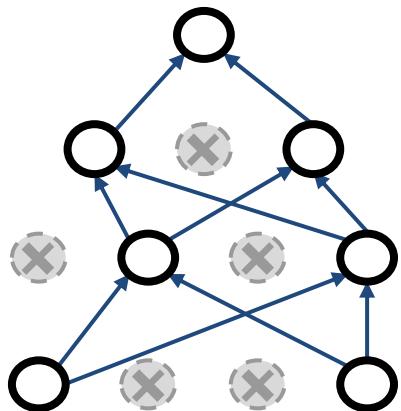
**Prevents co-adaptation of features**



# Regularization: Dropout

How can this possibly be a good idea?

- Another interpretation:
  - Dropout is training a large **ensemble** of models (that share parameters).

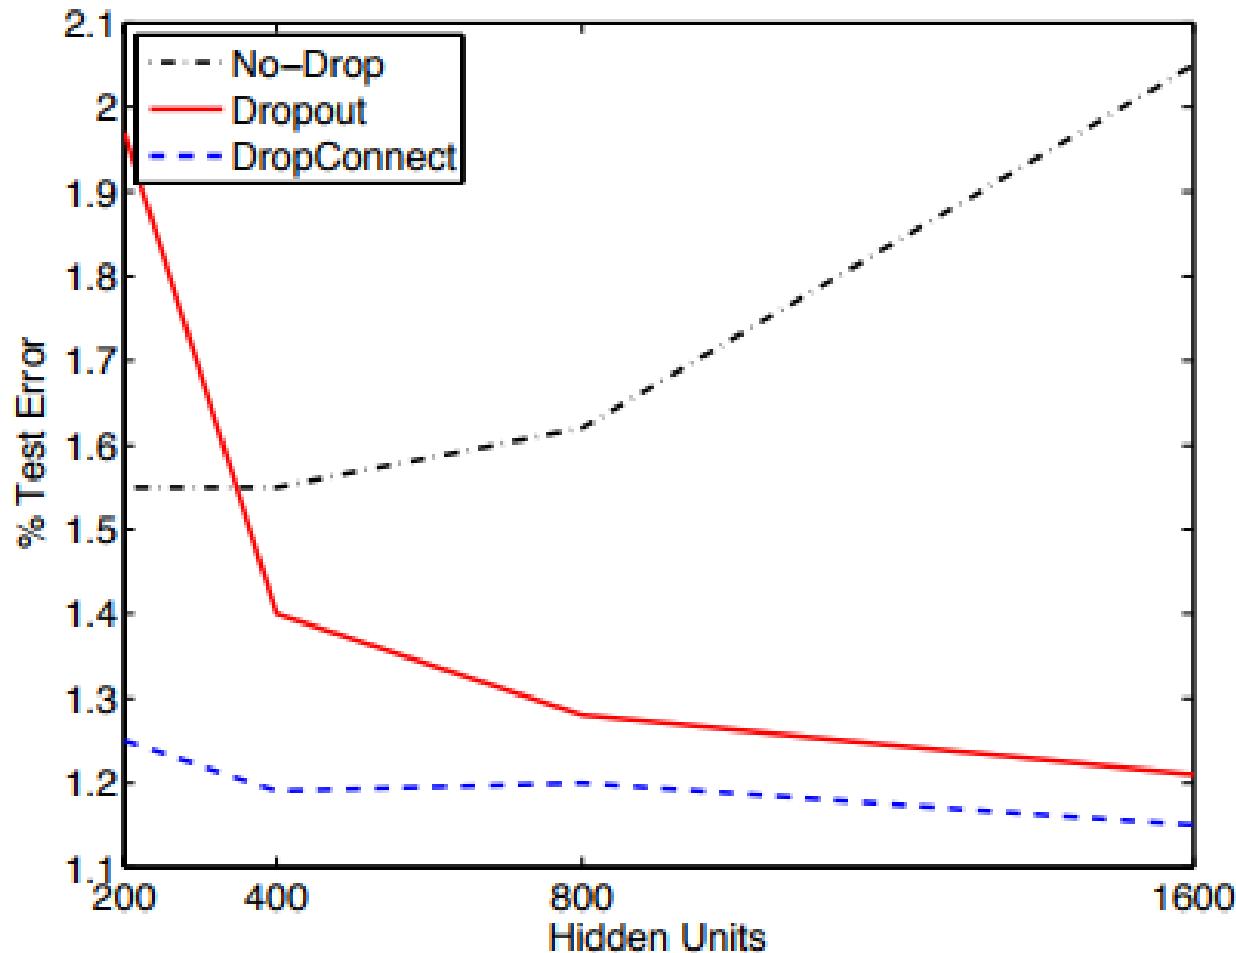


- Each binary mask is for one model
- An FC layer with 4096 units has

$2^{4096} \sim 10^{1233}$  possible masks!

Only  $\sim 10^{82}$  atoms in the universe...

# Drop-connection



# Internal Covariate Shifting

- Distribution of inputs to a **layer is changing during training**
- Difficult to train, solutions:
  - careful initialization
  - smaller learning rate
- Easier if **distribution of inputs stayed the same**
- How to **enforce same distribution?**

# Fighting with internal covariate shift

- Whitening would be a good first step
  - Would **remove nasty correlations**
- Problems with whitening?
  - Slow (have to do PCA for every layer)
  - **Cannot backprop through whitening**
- What is the best alternative?

# Batch Normalization(BN)

- Calculate basic statistics per batch to make mean  $\mu = 0$  and standard deviation  $\sigma = 1$
- **Doesn't eliminate correlations**
- Fast and **can backprop through it**
- How to compute the statistics?
  - Going over the entire dataset is too slow
  - Idea: the **batch is an approximation** of the dataset
  - Compute **statistics over the batch**

Mean  $\boxed{\mu_{\mathcal{B}}} \leftarrow \frac{1}{m} \sum_{i=1}^m \boxed{x_i}$  Batch size

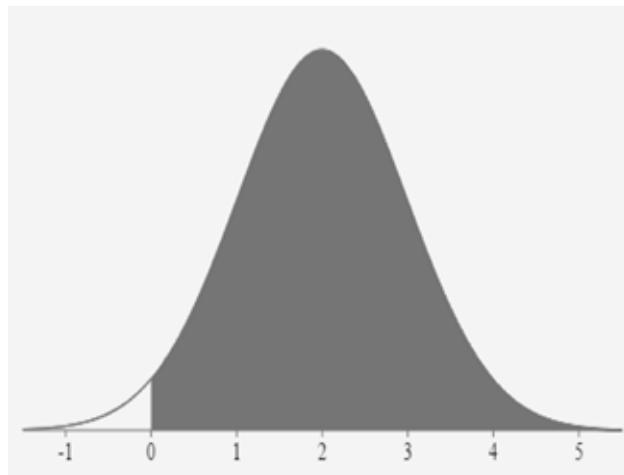
Variance  $\boxed{\sigma_{\mathcal{B}}^2} \leftarrow \frac{1}{m} \sum_{i=1}^m (\boxed{x_i} - \boxed{\mu_{\mathcal{B}}})^2$

Normalize  $\hat{x}_i \leftarrow \frac{\boxed{x_i} - \boxed{\mu_{\mathcal{B}}}}{\sqrt{\boxed{\sigma_{\mathcal{B}}^2} + \epsilon}}$

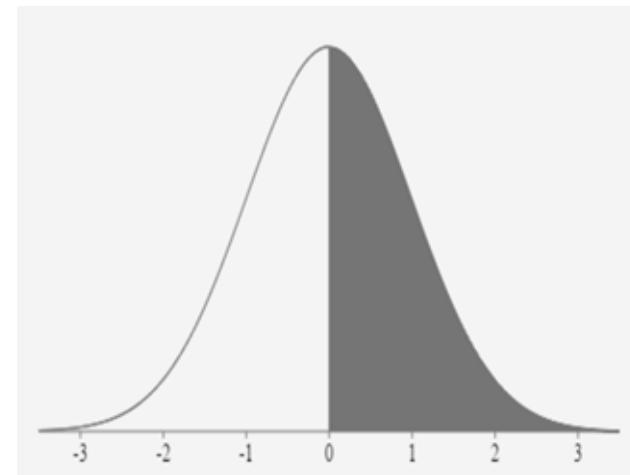
$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$$

$$y_i \leftarrow \boxed{\gamma} \hat{x}_i + \boxed{\beta} \equiv \text{BN}_{\gamma, \beta}(x_i)$$

Distribution of an activation before and after normalization



Before



After

# Summary

- Data (inputs) preprocessing
- Initialization, and activation functions
  - **One-sentence takeaway**
    - **Using compatible activations and gradient updates is crucial because deep learning is a repeated product of derivatives**
    - **if activations distort gradient scale, learning either vanishes or explodes, making efficient training impossible.**
- Regularization methods including:
  - Regularization methods
  - Batch normalization
- What's next?
  - Optimization methods and
  - Loss functions for training DNN models.

# Appendix

# Proof of He initialization

$$\mathbf{y}_l = W_l \mathbf{x}_l$$

$W$  is a  $d$ -by- $n$  matrix

$$\mathbf{y}_l = W_l \mathbf{x}_l$$

$W$  is a  $d$ -by- $n$  matrix

$$Var[y_l] = Var[w_l^1x_l^1 + \cdots + w_l^n x_l^n]$$

$$\mathbf{y}_l = W_l \mathbf{x}_l$$

$W$  is a  $d$ -by- $n$  matrix

$$Var[y_l] = Var[w_l^1 x_l^1 + \cdots + w_l^n x_l^n]$$

$w, x$  are both i.i.d. and independent of each other

$$\mathbf{y}_l = W_l \mathbf{x}_l$$

$W$  is a  $d$ -by- $n$  matrix

$$Var[y_l] = Var[w_l^1 x_l^1 + \cdots + w_l^n x_l^n]$$

$w, x$  are both i.i.d. and independent of each other

$$Var[y_l] = n_l Var[w_l x_l]$$

$$Var[y_l] = n_l Var[w_lx_l]$$

$$Var[y_l] = n_l Var[w_l x_l]$$

Let  $w_l$  have zero mean

$$Var[y_l] = n_l Var[w_l x_l]$$

Let  $w_l$  have zero mean

$$Var[ab] = E[a^2]E[b^2] + E[a]^2E[b]^2$$

$$Var[y_l] = n_l Var[w_l x_l]$$

Let  $w_l$  have zero mean

$$Var[ab] = E[a^2]E[b^2] + E[a]^2E[b]^2$$

$$Var[y_l] = n_l(E[w_l^2]E[x_l^2] + E[w_l]^2E[x_l]^2)$$

$$Var[y_l] = n_l Var[w_l x_l]$$

Let  $w_l$  have zero mean

$$Var[ab] = E[a^2]E[b^2] + E[a]^2E[b]^2$$

$$Var[y_l] = n_l(E[w_l^2]E[x_l^2] + E[w_l]^2E[x_l]^2)$$

$$Var[y_l] = n_l E[w_l^2] E[x_l^2]$$

$$Var[y_l] = n_l E[w_l^2]E[x_l^2]$$

$$Var[a]=E[(a-E[a])^2]$$

$$Var[y_l] = n_l E[w_l^2]E[x_l^2]$$

$$Var[a]=E[(a-E[a])^2]$$

$$Var[y_l] = n_l Var[w_l]E[x_l^2]$$

$$Var[y_l] = n_l Var[w_l] E[x_l^2]$$

$$x_l = \max(0, y_{l-1})$$

Let  $w_{l-1}$  have a zero-mean symmetric distribution

$$Var[y_l] = n_l Var[w_l] E[x_l^2]$$

$$x_l = \max(0, y_{l-1})$$

Let  $w_{l-1}$  have a zero-mean symmetric distribution

Then  $y_{l-1}$  also has a zero-mean symmetric distribution

$$Var[y_l] = n_l Var[w_l] E[x_l^2]$$

$$x_l = \max(0, y_{l-1})$$

Let  $w_{l-1}$  have a zero-mean symmetric distribution

Then  $y_{l-1}$  also has a zero-mean symmetric distribution

$$Var[y_{l-1}] = \int_{-\infty}^0 y_{l-1}^2 p(y_{l-1}) dy_{l-1} + \int_0^\infty y_{l-1}^2 p(y_{l-1}) dy_{l-1}$$

$$Var[y_l] = n_l Var[w_l] E[x_l^2]$$

$$x_l = \max(0, y_{l-1})$$

Let  $w_{l-1}$  have a zero-mean symmetric distribution

Then  $y_{l-1}$  also has a zero-mean symmetric distribution

$$Var[y_{l-1}] = \int_{-\infty}^0 y_{l-1}^2 p(y_{l-1}) dy_{l-1} + \int_0^\infty y_{l-1}^2 p(y_{l-1}) dy_{l-1}$$

$$Var[y_{l-1}] = 2 \int_0^\infty y_{l-1}^2 p(y_{l-1}) dy_{l-1}$$

$$Var[y_{l-1}] = 2\int_0^{\infty}y_{l-1}^2p(y_{l-1})dy_{l-1}$$

$$Var[y_{l-1}] = 2 \int_0^{\infty} y_{l-1}^2 p(y_{l-1}) dy_{l-1}$$

$$Var[y_{l-1}] = 2 \int_0^{\infty} x_l^2 p(x_l) dx_l$$

$$Var[y_{l-1}] = 2 \int_0^{\infty} y_{l-1}^2 p(y_{l-1}) dy_{l-1}$$

$$Var[y_{l-1}] = 2 \int_0^{\infty} x_l^2 p(x_l) dx_l$$

$$\frac{1}{2}Var[y_{l-1}] = E[x_l^2]$$

$$Var[y_{l-1}] = 2 \int_0^{\infty} y_{l-1}^2 p(y_{l-1}) dy_{l-1}$$

$$Var[y_{l-1}] = 2 \int_0^{\infty} x_l^2 p(x_l) dx_l$$

$$\frac{1}{2}Var[y_{l-1}] = E[x_l^2]$$

$$Var[y_l] = n_l Var[w_l] E[x_l^2]$$

$$Var[y_{l-1}] = 2 \int_0^{\infty} y_{l-1}^2 p(y_{l-1}) dy_{l-1}$$

$$Var[y_{l-1}] = 2 \int_0^{\infty} x_l^2 p(x_l) dx_l$$

$$\frac{1}{2}Var[y_{l-1}] = E[x_l^2]$$

$$Var[y_l] = n_l Var[w_l] E[x_l^2]$$

$$Var[y_l] = \frac{1}{2}n_l Var[w_l] Var[y_{l-1}]$$

$$Var[y_l] = \frac{1}{2} n_l Var[w_l] Var[y_{l-1}]$$

$$Var[y_l] = \frac{1}{2} n_l Var[w_l] Var[y_{l-1}]$$

$$Var[y_L] = Var[y_1]\prod_{l=2}^L\left(\frac{1}{2}n_lVar[w_l]\right)$$

$$Var[y_l] = \frac{1}{2} n_l Var[w_l] Var[y_{l-1}]$$

$$Var[y_L] = Var[y_1]\prod_{l=2}^L\left(\frac{1}{2}n_lVar[w_l]\right)$$

$$Var[y_L]=Var[y_1]$$

$$Var[y_l] = \frac{1}{2} n_l Var[w_l] Var[y_{l-1}]$$

$$Var[y_L] = Var[y_1]\prod_{l=2}^L\left(\frac{1}{2}n_lVar[w_l]\right)$$

$$Var[y_L]=Var[y_1]$$

$$\frac{1}{2}n_lVar[w_l]=1$$

$$\frac{1}{2}n_lVar[w_l]=1$$

$$\frac{1}{2}n_lVar[w_l]=1$$

$$w_l \sim \mathcal{N}\left(0,\frac{2}{n_l}\right)$$