

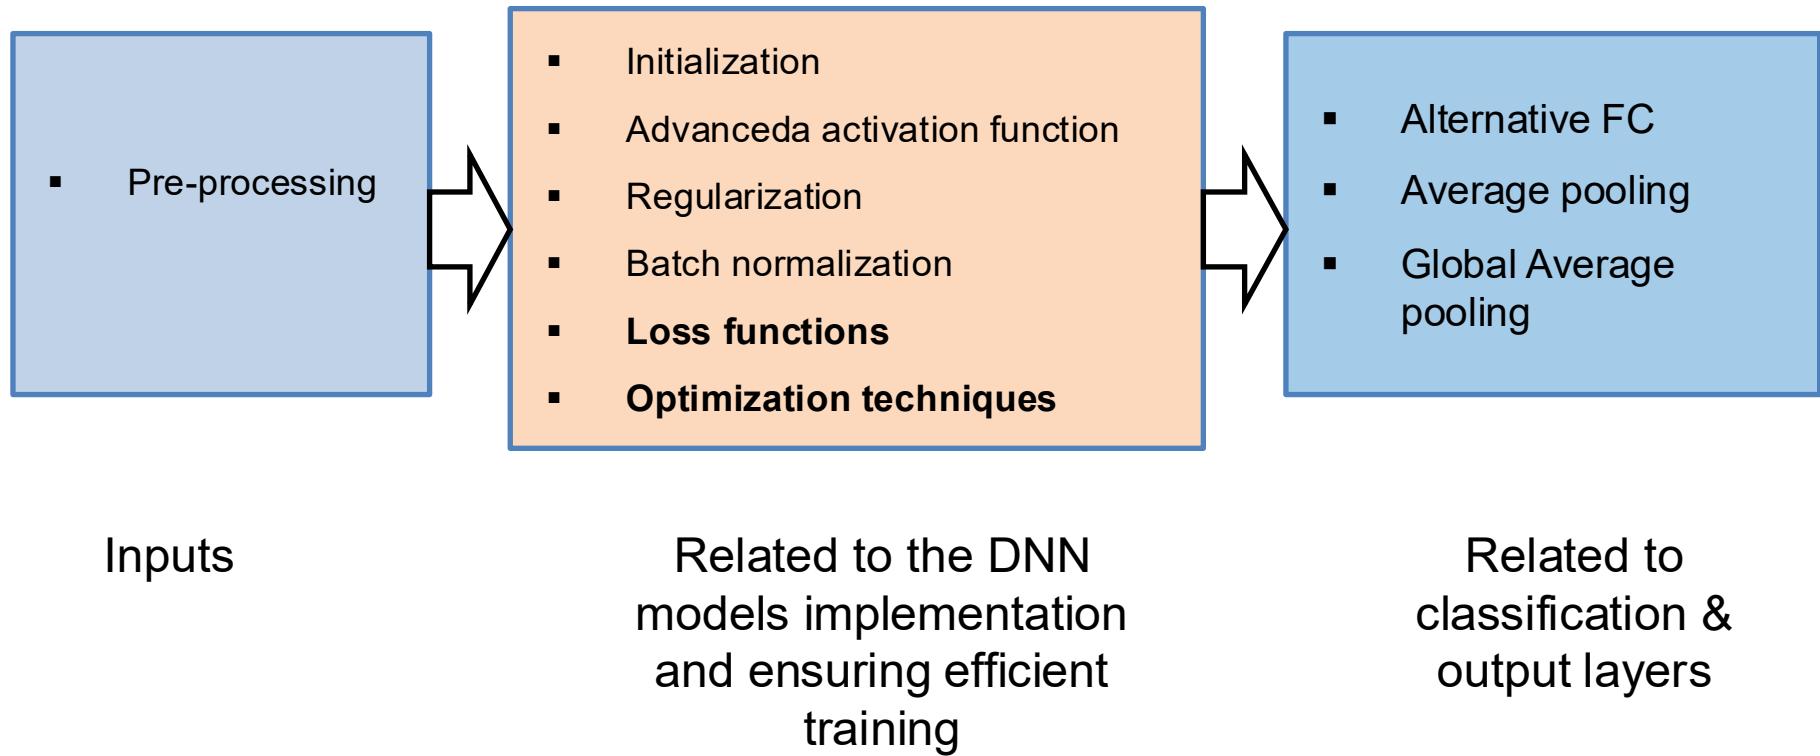
COMP/EECE 7/8740 Neural Networks

Topics: Advanced Training Techniques

- Loss functions
 - Zero-one loss (per class)
 - Binary and Categorical Cross Entropy Loss
 - Focal Loss
- Optimization techniques:
 - SGD, Adagrad, AdaDelta, RMSProp, Adam, and AdamW
 - Sophia and Muon losses

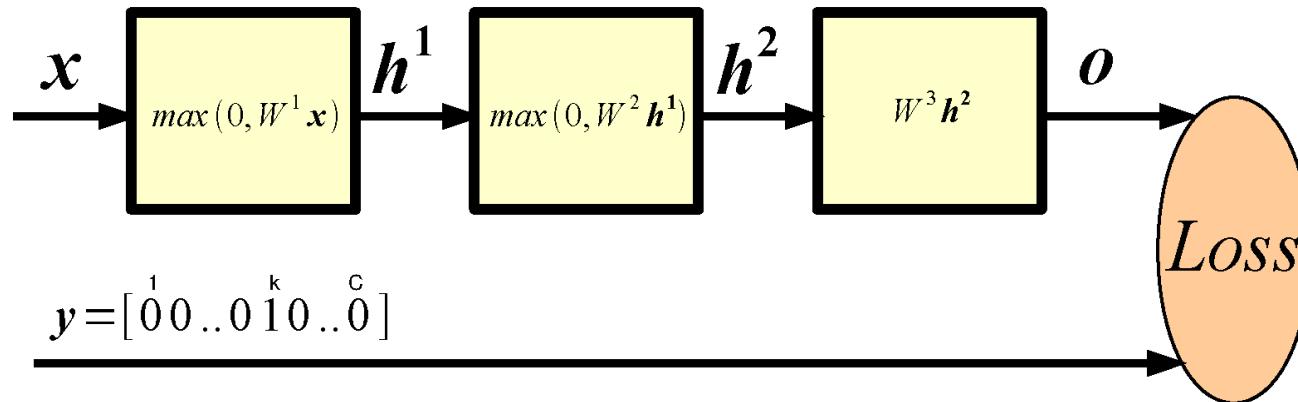
Md Zahangir Alom
Department of Computer Science
University of Memphis, TN

Overview : Learning with DNN



Loss Functions

How Good is a Network?



- A **loss function** measures how far a model's predictions are from the true targets
- Training aims to **minimize this loss**, guiding parameter updates so the model makes more accurate predictions.

What is an appropriate loss?

- Compare training class to output class
- Zero-one loss (per class)

$$L(\hat{y}, y) = I(\hat{y} \neq y),$$

- **Is it good?**
 - Nope – **it's a step function.**
 - We need to compute the gradient of the loss.
 - This loss is **not differentiable**, and ‘flips’ easily.

Entropy Loss function

$$P(y | x) = \begin{cases} \hat{p} & \text{if } y = 1 \\ 1 - \hat{p} & \text{if } y = 0 \end{cases}$$

Max/Min occurs at points where derivative is equal to 0:

$$\frac{\partial P(y | x)}{\partial p} = 0$$

Our examples can only be part of one class at a time either 1 or 0 so,

- when $y = 1$ then \hat{p} should be the output and $1 - \hat{p}$ should be ignored,
- similarly, when $y = 0$ then $1 - \hat{p}$ should be the output and \hat{p} should be ignored.

We can combine this as follows:

$$P(y | x) = \hat{p}^y(1 - \hat{p})^{1-y}$$

if $y = 1$ then:

$$P(y | x) = \hat{p}^y(1 - \hat{p})^{1-y} = \hat{p}^1(1 - \hat{p})^{1-1} = \hat{p}$$

if $y = 0$ then:

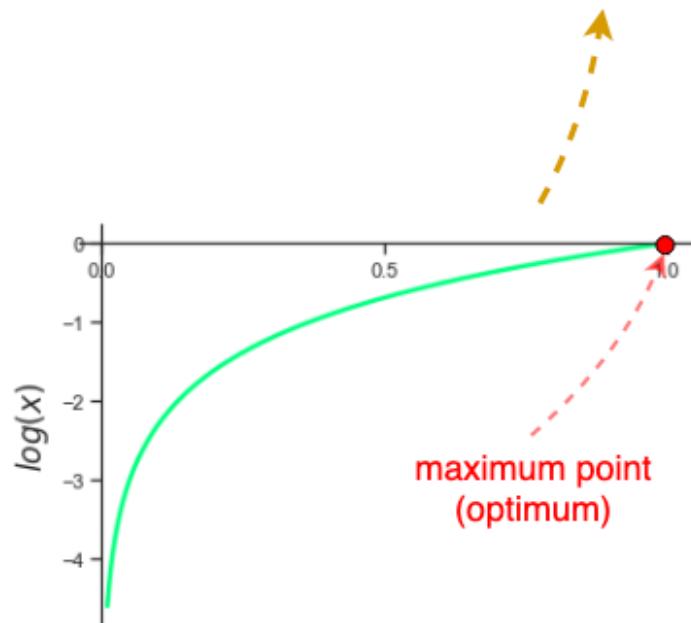
$$P(y | x) = \hat{p}^y(1 - \hat{p})^{1-y} = \hat{p}^0(1 - \hat{p})^{1-0} = 1 - \hat{p}$$

Entropy Loss function

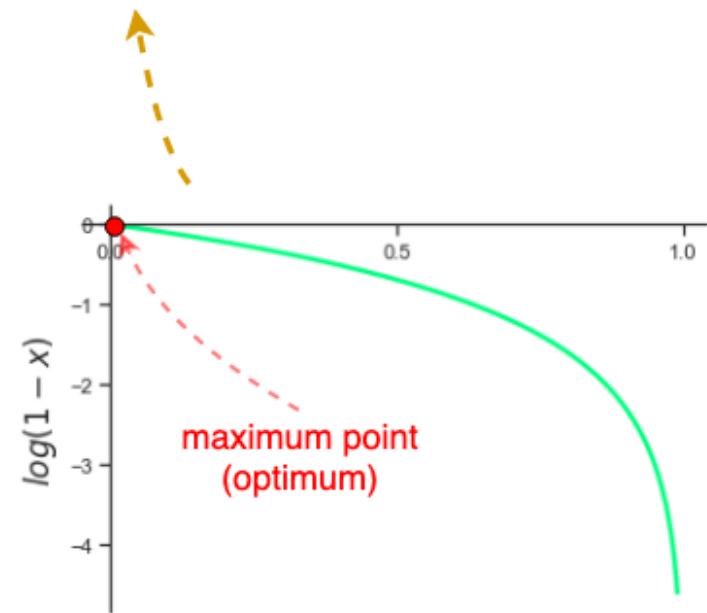
Recall the properties of natural log:

- $\log(a * b) = \log(a) + \log(b)$
- $\log(a^x) = x \log(a)$

$$\begin{aligned}\log(P(y | x)) &= \log(\hat{p}^y(1 - \hat{p})^{(1-y)}) \\ &= \log(\hat{p}^y) + \log((1 - \hat{p})^{(1-y)}) \\ &= y \log(\hat{p}) + (1 - y) \log(1 - \hat{p})\end{aligned}$$



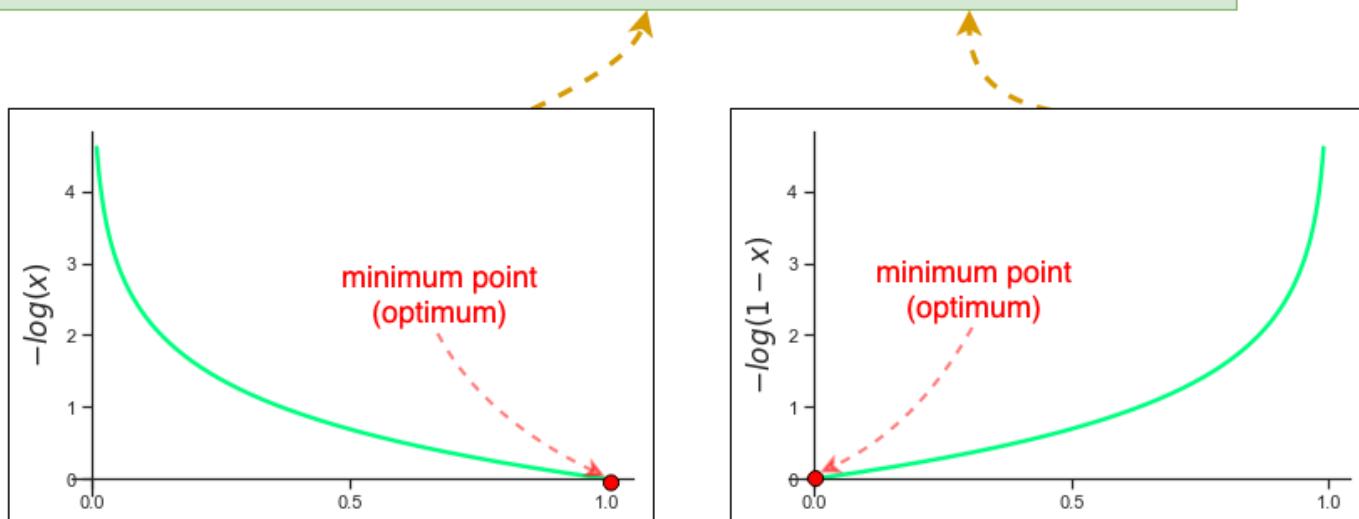
maximum point
(optimum)



maximum point
(optimum)

Entropy Loss function

$$\begin{aligned} \text{Binary Cross Entropy Loss}(y, \hat{p}) &= -\log(P(y | x)) \\ &= -(y \log(\hat{p}) + (1 - y) \log(1 - \hat{p})) \\ &= -y \log(\hat{p}) - (1 - y) \log(1 - \hat{p}) \end{aligned}$$



Binary Cross Entropy Loss/Cost function

$$\begin{aligned}\text{Binary Cross-Entropy Cost} &= \frac{1}{m} \left(-\log \left(\mathcal{L}(\hat{p}^{(1)}, \hat{p}^{(2)}, \dots, \hat{p}^{(m)}) \right) \right) \\ &= \frac{1}{m} \sum_{i=1}^m -y^{(i)} \log(\hat{p}^{(i)}) - (1 - y^{(i)}) \log(1 - \hat{p}^{(i)})\end{aligned}$$

$$\begin{aligned}-\log \left(\mathcal{L}(\hat{p}^{(1)}, \hat{p}^{(2)}, \dots, \hat{p}^{(m)}) \right) &= -\sum_{i=1}^m y^{(i)} \log(\hat{p}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{p}^{(i)}) \\ &= \sum_{i=1}^m -y^{(i)} \log(\hat{p}^{(i)}) - (1 - y^{(i)}) \log(1 - \hat{p}^{(i)})\end{aligned}$$

Categorical Cross Entropy Loss/Cost function

$$\mathcal{L}(\hat{p}^{(1)}, \hat{p}^{(2)}, \dots, \hat{p}^{(m)} \mid (x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})) = \prod_{i=1}^m \hat{p}^{(i)y^{(i)}} (1 - \hat{p}^{(i)})^{1-y^{(i)}}$$

This is the ***likelihood*(\mathcal{L})** function, where we are trying to maximize the each probability, \hat{p} , given the examples $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})$.

$\prod_{i=1}^m$ represents product over examples 1 to m

$$\begin{aligned}\log \left(\mathcal{L}(\hat{p}^{(1)}, \hat{p}^{(2)}, \dots, \hat{p}^{(m)} \mid (x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})) \right) &= \log \left(\prod_{i=1}^m \hat{p}^{(i)y^{(i)}} (1 - \hat{p}^{(i)})^{1-y^{(i)}} \right) \\ &= \sum_{i=1}^m \log \left(\hat{p}^{(i)y^{(i)}} (1 - \hat{p}^{(i)})^{1-y^{(i)}} \right) \\ &= \sum_{i=1}^m y^{(i)} \log(\hat{p}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{p}^{(i)})\end{aligned}$$

This is called the ***log-likelihood*** function

Why cross-entropy loss function?

Entropy :

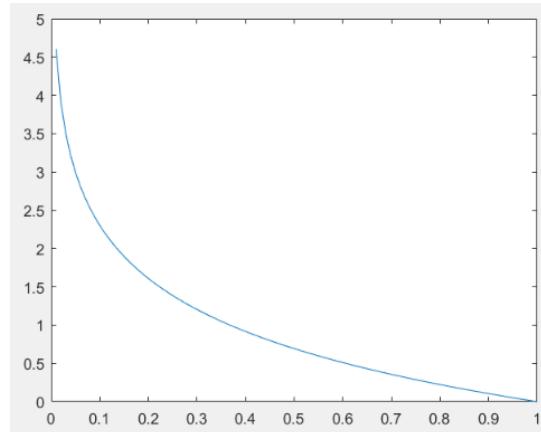
$$H(X) = \begin{cases} - \int_x p(x) \log p(x), & \text{if } X \text{ is continuous} \\ - \sum_x p(x) \log p(x), & \text{if } X \text{ is discrete} \end{cases}$$

Cross-entropy is defined as

$$L_{\text{CE}} = - \sum_{i=1}^n t_i \log(p_i), \text{ for n classes,}$$

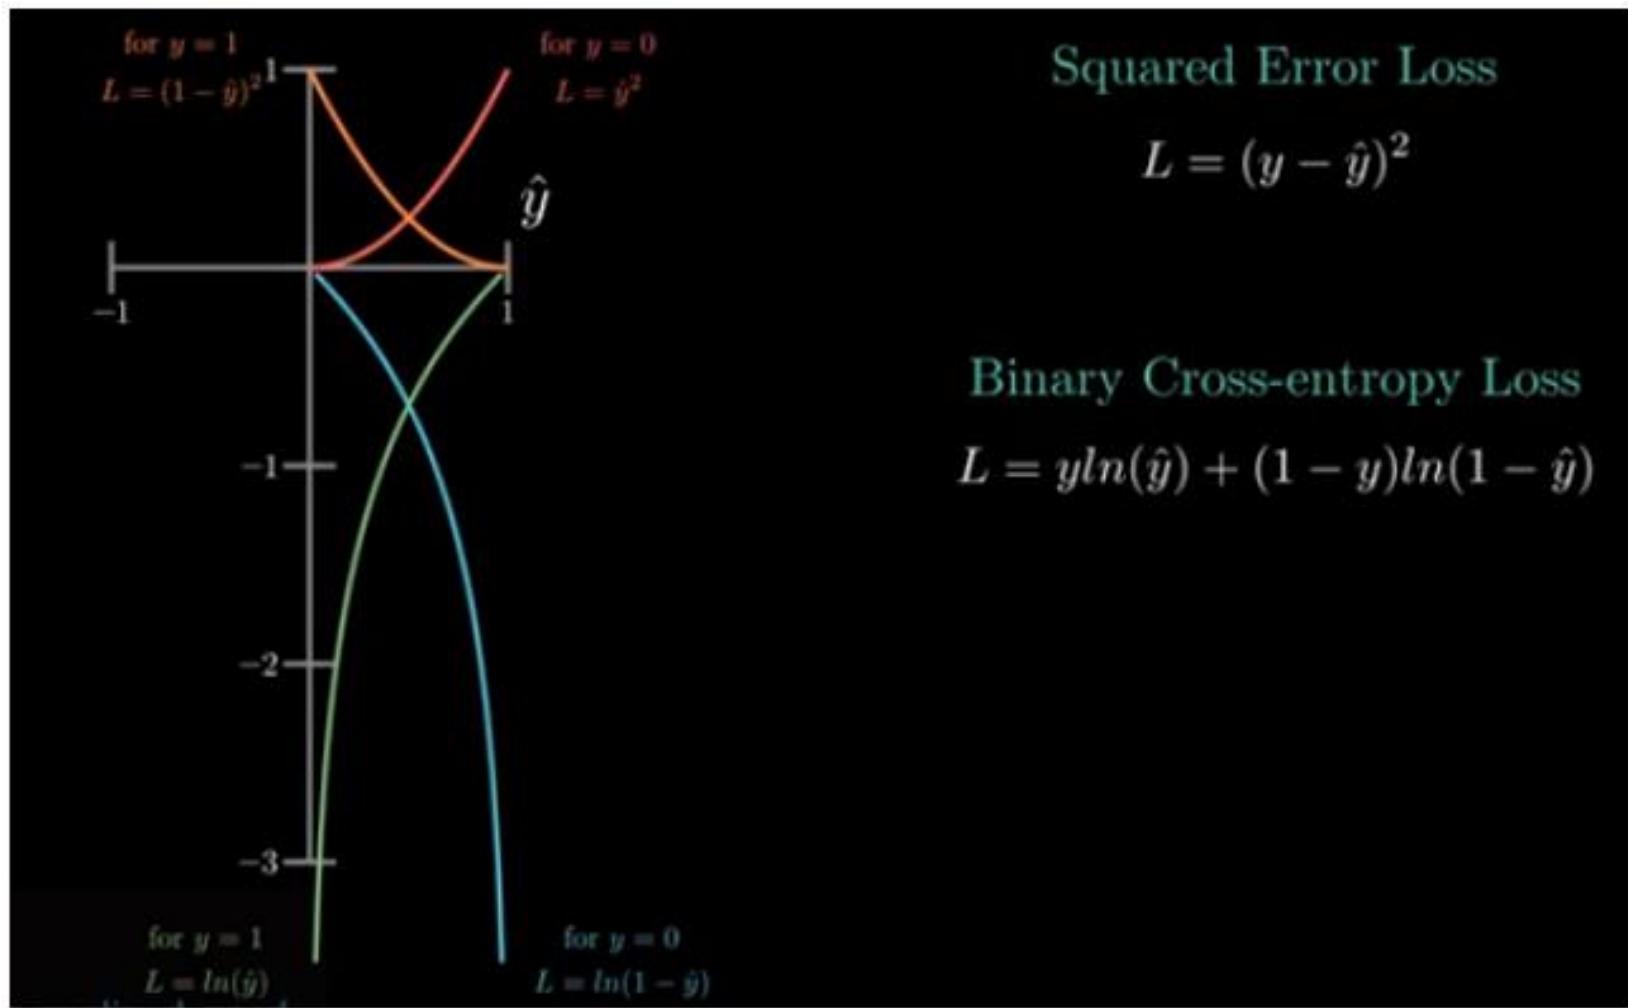
where t_i is the truth label and p_i is the Softmax probability for the i^{th} class.

- **Is it a good loss?**
 - Differentiable
 - Cost decreases as probability increases



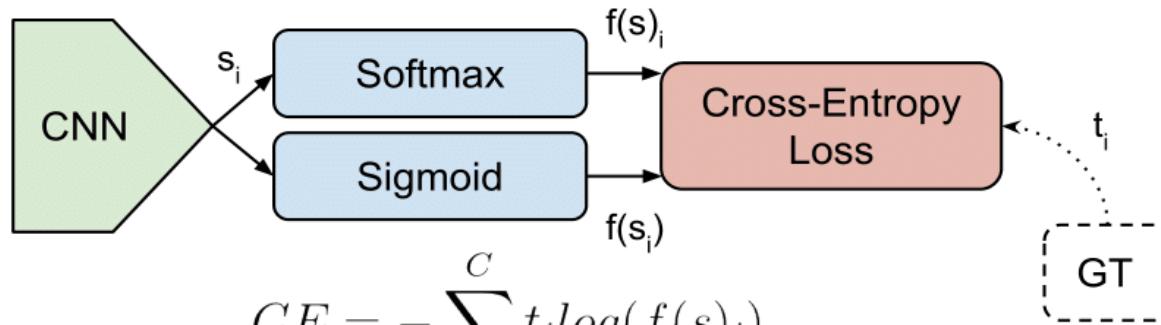
Why cross-entropy loss function?

Gradient Scope in Loss functions



Types of Cross-Entropy Loss

- Categorical Cross-Entropy Loss
- Binary Cross-Entropy Loss and

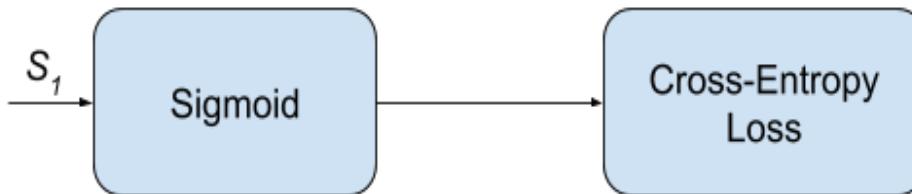


$$CE = - \sum_i^C t_i \log(f(s)_i)$$

$$CE = - \sum_{i=1}^{C'=2} t_i \log(f(s_i)) = -t_1 \log(f(s_1)) - (1 - t_1) \log(1 - f(s_1))$$

Binary Cross-Entropy Loss

- Also called **Sigmoid Cross-Entropy loss**. It is a **Sigmoid activation** plus a **Cross-Entropy loss**.
- It's called **Binary Cross-Entropy Loss** because it sets up a binary classification problem between $C'=2$ classes for every class in C.



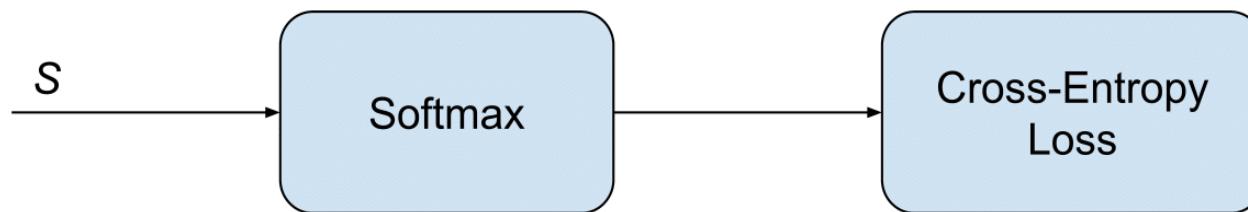
$$CE = -t_1 \log(f(s_1)) - (1 - t_1) \log(1 - f(s_1))$$

$$\begin{aligned} L &= -\sum_{i=1}^2 t_i \log(p_i) \\ &= -[t_1 \log(p_1) + t_2 \log(p_2)] \\ &= -[t \log(p) + (1 - t) \log(1 - p)] \end{aligned}$$

$$P = f(s_i) = \frac{1}{1 + e^{-s_i}}$$

Categorical Cross-Entropy Loss

- Also called **Softmax Loss**. It is a **Softmax activation** plus a **Cross-Entropy loss**
- We use this loss when we train a DNN or CNN to output a probability over the C classes for each image (multi-class classification).

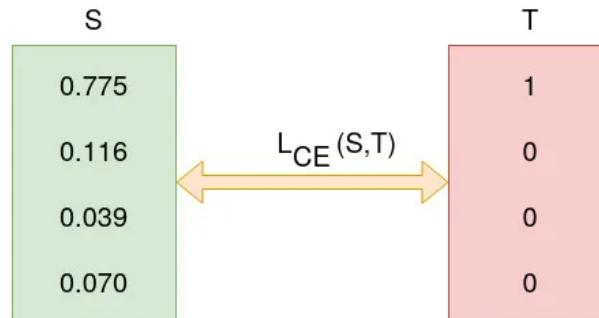


$$f(s)_i = \frac{e^{s_i}}{\sum_j^C e^{s_j}} \quad CE = - \sum_i^C t_i \log(f(s)_i)$$

Examples

Categorical Cross-Entropy Loss

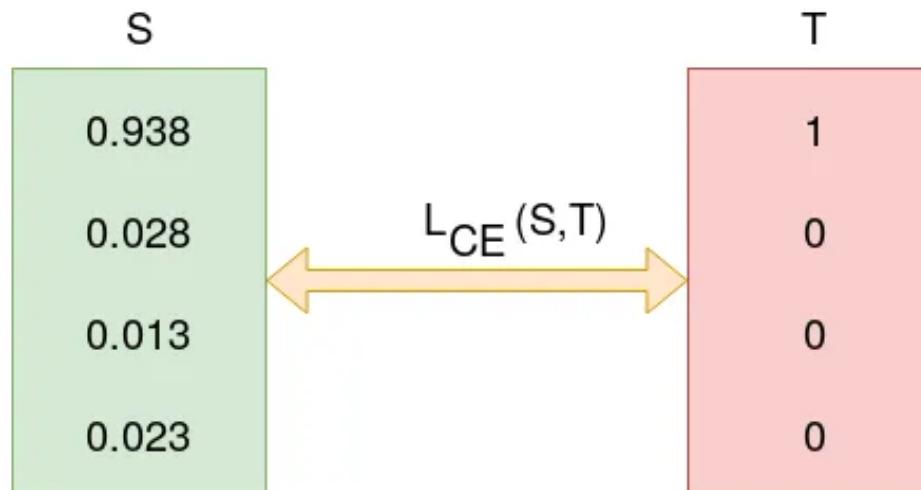
$$L = -\frac{1}{N} \left[\sum_{j=1}^N [t_j \log(p_j) + (1 - t_j) \log(1 - p_j)] \right]$$



$$\begin{aligned} L_{CE} &= - \sum_{i=1} T_i \log(S_i) \\ &= - [1 \log_2(0.775) + 0 \log_2(0.126) + 0 \log_2(0.039) + 0 \log_2(0.070)] \\ &= - \log_2(0.775) \\ &= 0.3677 \end{aligned}$$

Examples:

Categorical Cross-Entropy Loss



$$\begin{aligned}L_{CE} &= -1 \log_2(0.936) + 0 + 0 + 0 \\&= 0.095\end{aligned}$$

Softmax Classifier (Multinomial Logistic Regression)

- Softmax layer as output layer

Ordinary Layer

$$z_1 \rightarrow \sigma \rightarrow y_1 = \sigma(z_1)$$

$$z_2 \rightarrow \sigma \rightarrow y_2 = \sigma(z_2)$$

$$z_3 \rightarrow \sigma \rightarrow y_3 = \sigma(z_3)$$

- In general, the output of network can be any value.
- May not be easy to interpret

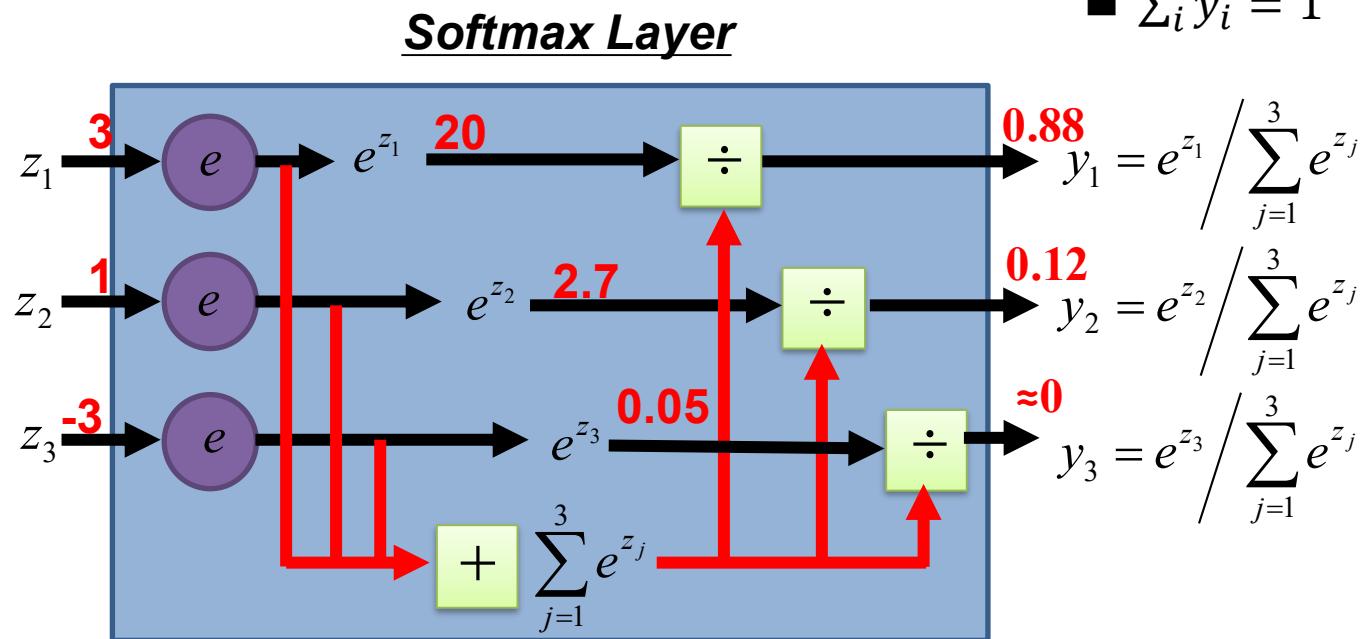
Softmax Classifier (Multinomial Logistic Regression)

- Softmax layer as the output layer

Probability:

$$\blacksquare 1 > y_i > 0$$

$$\blacksquare \sum_i y_i = 1$$



Softmax Classifier (Multinomial Logistic Regression)



Want to interpret raw classifier scores as **probabilities**

| | |
|------|------------|
| cat | 3.2 |
| car | 5.1 |
| frog | -1.7 |

Softmax Classifier (Multinomial Logistic Regression)



Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

| | |
|------|------------|
| cat | 3.2 |
| car | 5.1 |
| frog | -1.7 |

Softmax Classifier (Multinomial Logistic Regression)

Want to interpret raw classifier scores as **probabilities**



$$s = f(x_i; W)$$

$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

Probabilities
must be ≥ 0

| | |
|------|------|
| cat | 3.2 |
| car | 5.1 |
| frog | -1.7 |

\exp

| |
|-------|
| 24.5 |
| 164.0 |
| 0.18 |

unnormalized
probabilities

Softmax Classifier (Multinomial Logistic Regression)



Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

| | |
|------|------|
| cat | 3.2 |
| car | 5.1 |
| frog | -1.7 |

Probabilities
must be ≥ 0

$\xrightarrow{\text{exp}}$

| |
|-------|
| 24.5 |
| 164.0 |
| 0.18 |

unnormalized
probabilities

Probabilities
must sum to 1

$\xrightarrow{\text{normalize}}$

| |
|------|
| 0.13 |
| 0.87 |
| 0.00 |

probabilities

Softmax Classifier (Multinomial Logistic Regression)



Want to interpret raw classifier scores as **probabilities**

$$\mathbf{s} = f(\mathbf{x}_i; \mathbf{W})$$

$$P(Y = k | X = \mathbf{x}_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

cat
car
frog

| |
|------|
| 3.2 |
| 5.1 |
| -1.7 |

Unnormalized log-probabilities / logits

Probabilities must be ≥ 0

| |
|-------|
| 24.5 |
| 164.0 |
| 0.18 |

unnormalized probabilities

Probabilities must sum to 1

| |
|------|
| 0.13 |
| 0.87 |
| 0.00 |

probabilities

Softmax Classifier (Multinomial Logistic Regression)



Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

$$L_i = -\log P(Y = y_i|X = x_i)$$

cat
car
frog

| |
|------|
| 3.2 |
| 5.1 |
| -1.7 |

Unnormalized log-probabilities / logits

Probabilities must be ≥ 0

| |
|-------|
| 24.5 |
| 164.0 |
| 0.18 |

unnormalized probabilities

Probabilities must sum to 1

| |
|------|
| 0.13 |
| 0.87 |
| 0.00 |

probabilities

exp

normalize

$$\rightarrow L_i = -\log(0.13) = 0.89$$

Softmax Classifier (Multinomial Logistic Regression)



Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

$$L_i = -\log P(Y = y_i|X = x_i)$$

Probabilities
must be ≥ 0

Probabilities
must sum to 1

cat
car
frog

| |
|------|
| 3.2 |
| 5.1 |
| -1.7 |

Unnormalized log-
probabilities / logits

| |
|-------|
| 24.5 |
| 164.0 |
| 0.18 |

unnormalized
probabilities

exp

normalize

| |
|------|
| 0.13 |
| 0.87 |
| 0.00 |

probabilities

$$\rightarrow L_i = -\log(0.87) \\ = 0.0604$$

Maximum Likelihood
Estimation :
Choose probabilities to
maximize the likelihood of
the observed data

Softmax Classifier (Multinomial Logistic Regression)



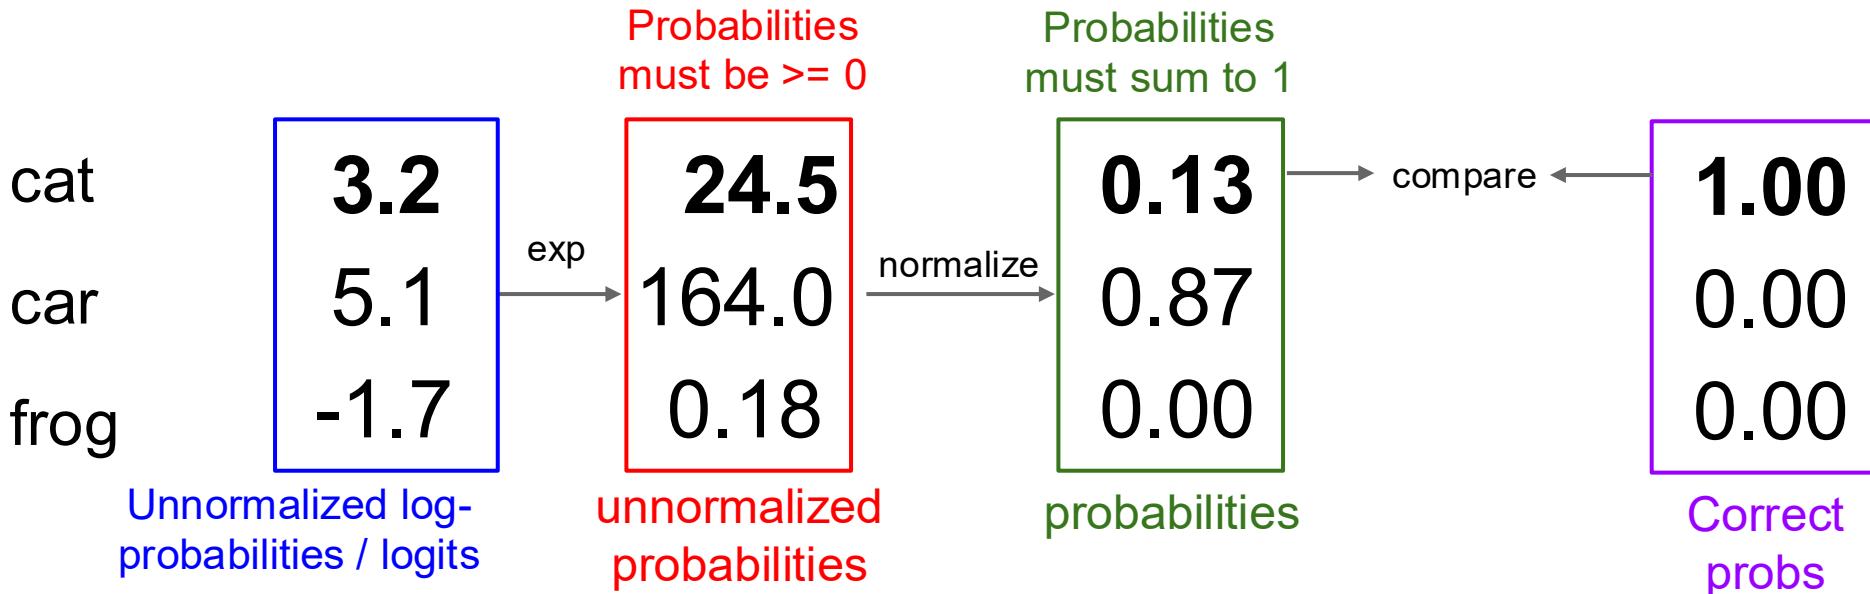
Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

$$L_i = -\log P(Y = y_i | X = x_i)$$



Softmax Classifier (Multinomial Logistic Regression)



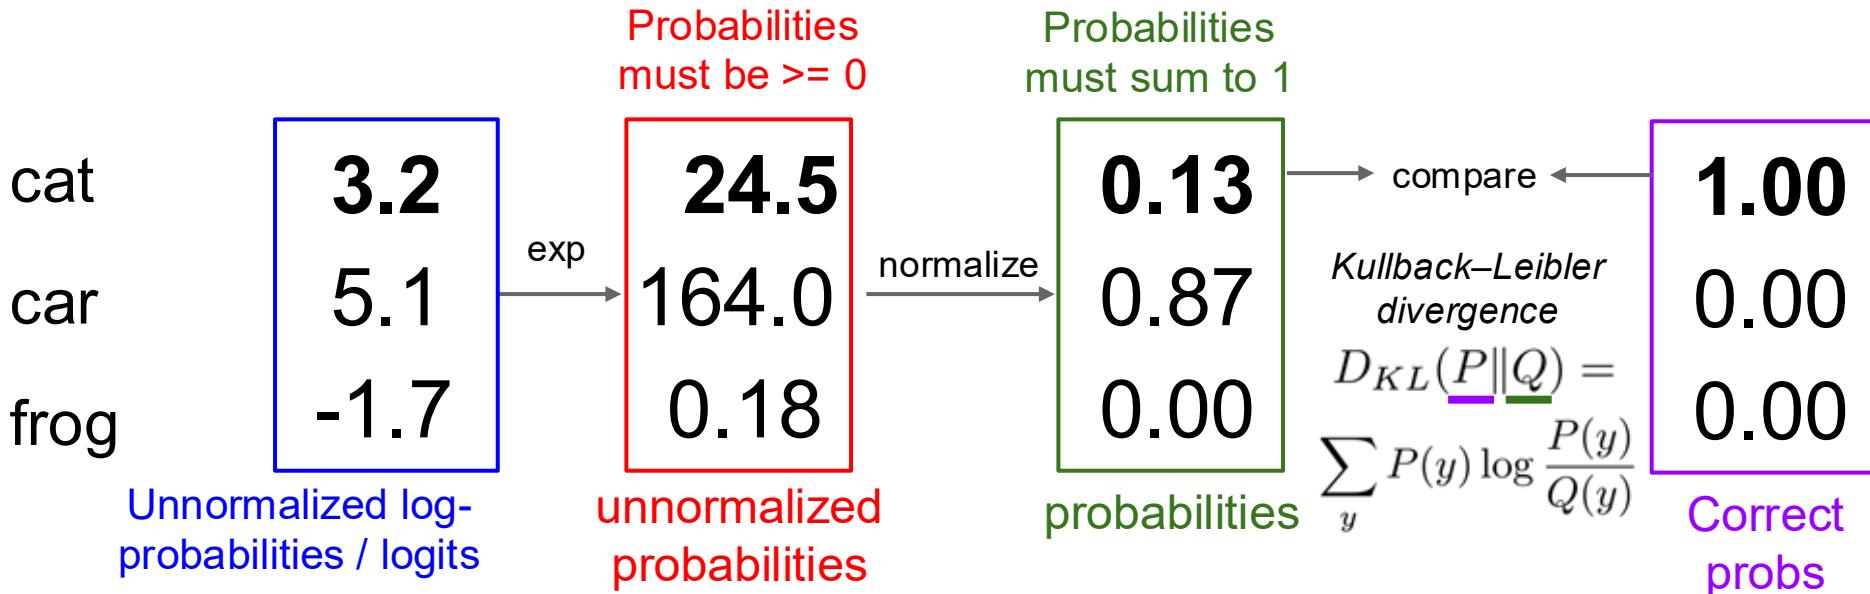
Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

$$L_i = -\log P(Y = y_i | X = x_i)$$



Softmax Classifier (Multinomial Logistic Regression)



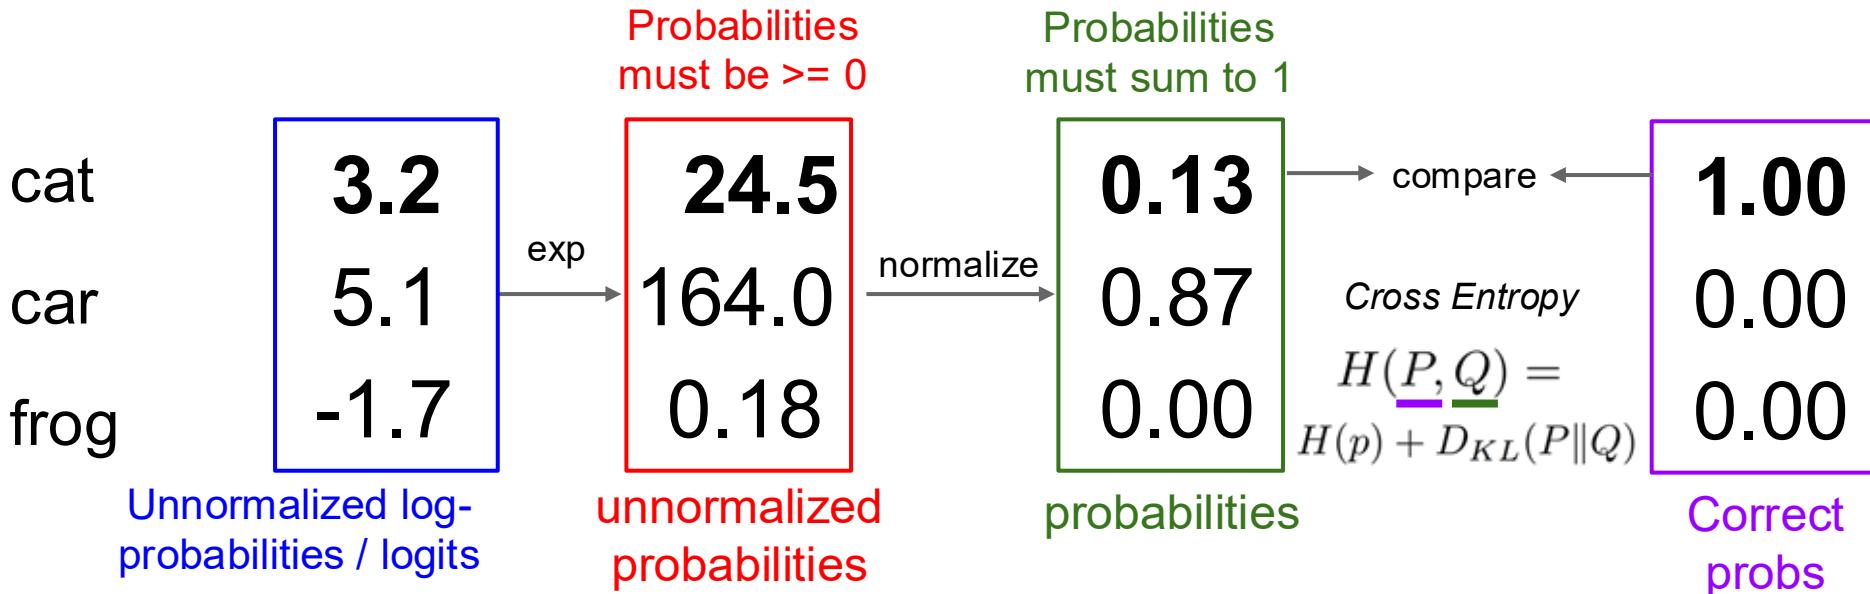
Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

$$L_i = -\log P(Y = y_i | X = x_i)$$



Softmax Classifier (Multinomial Logistic Regression)



Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax
Function

Maximize probability of correct class

$$L_i = -\log P(Y = y_i | X = x_i)$$

Putting it all together:

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

| | |
|------|-------------|
| cat | 3.2 |
| car | 5.1 |
| frog | -1.7 |

Softmax Classifier (Multinomial Logistic Regression)



Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

Maximize probability of correct class

$$L_i = -\log P(Y = y_i|X = x_i)$$

Putting it all together:

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

| | |
|------|------------|
| cat | 3.2 |
| car | 5.1 |
| frog | -1.7 |

Q: What is the min/max possible loss L_i ?

Softmax Classifier (Multinomial Logistic Regression)



| | |
|------|------|
| cat | 3.2 |
| car | 5.1 |
| frog | -1.7 |

Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

Maximize probability of correct class

$$L_i = -\log P(Y = y_i|X = x_i)$$

Putting it all together:

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

Q: What is the min/max possible loss L_i ?

A: min 0, max infinity

Softmax Classifier (Multinomial Logistic Regression)

| | |
|------|------|
| cat | 3.2 |
| car | 5.1 |
| frog | -1.7 |



Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax Function

Maximize probability of correct class

$$L_i = -\log P(Y = y_i | X = x_i)$$

Putting it all together:

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

Q2: At initialization all s will be approximately equal; what is the loss?

Softmax Classifier (Multinomial Logistic Regression)



Want to interpret raw classifier scores as **probabilities**

$$s = f(x_i; W)$$

$$P(Y = k | X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Softmax
Function

Maximize probability of correct class

$$L_i = -\log P(Y = y_i | X = x_i)$$

Putting it all together:

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

| | |
|------|------|
| cat | 3.2 |
| car | 5.1 |
| frog | -1.7 |

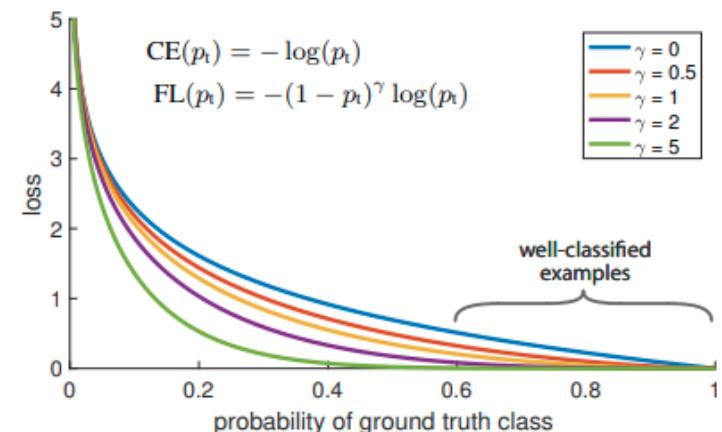
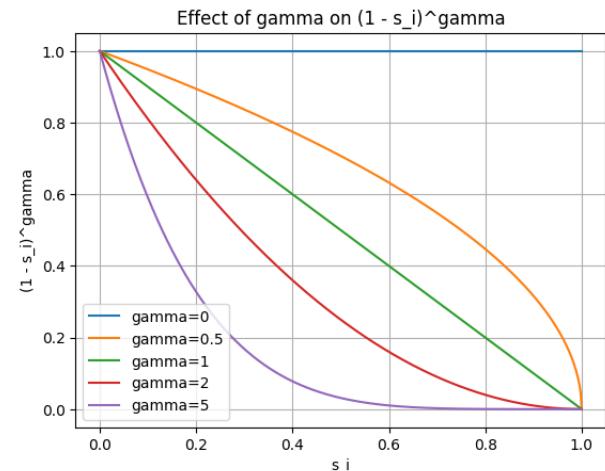
Q2: At initialization all s will be approximately equal; what is the loss?
A: $\log(C)$, eg $\log(10) \approx 2.3$

Focal Loss

- **Focal Loss** was introduced by Lin et al., from Facebook. They claim to improve one-stage object detectors using **Focal Loss** to train a detector named “**RetinaNet**”.
- **Focal loss** is a **Cross-Entropy Loss** that weights the contribution of each sample to the loss based in the classification error.
- **Focal loss** could also be considered a **Binary Cross-Entropy Loss** (**Sigmoid activations + Cross-Entropy Loss**)

$$FL = - \sum_{i=1}^{C=2} (1 - s_i)^\gamma t_i \log(s_i) \quad \gamma >= 0$$

$\gamma=0$, Focal Loss is equivalent to Binary Cross Entropy Loss.



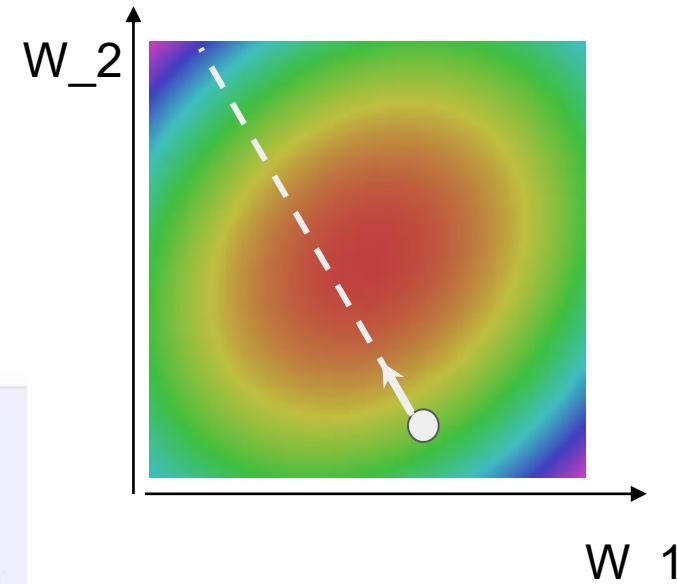
DNN optimization methods

DNN optimization methods are algorithms (like SGD, Adam, or AdamW) that **iteratively update a neural network's parameters to minimize the loss function** during training.

- Stochastic Gradient Descent (SGD)
- SGD+Momentum
- Adagrad
- Nesterov's Accelerated Gradient (NAG)
- RMSProp
- Adam
- AdamW and
- Many more ...

```
# Vanilla Gradient Descent
```

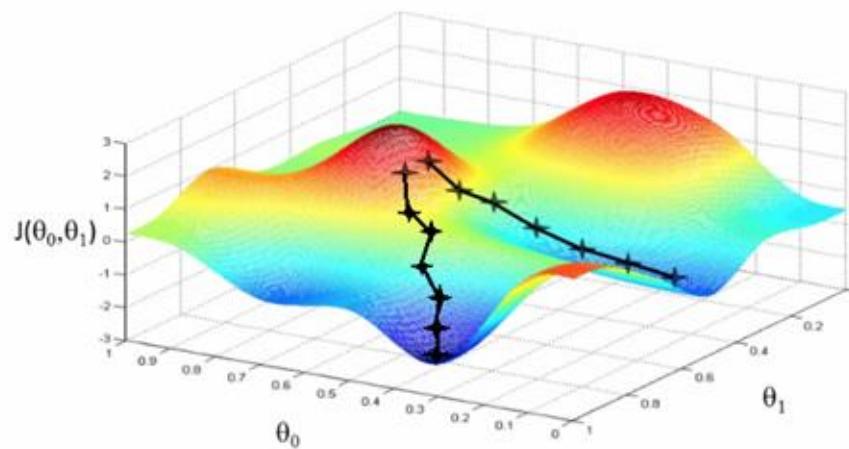
```
while True:  
    weights_grad = evaluate_gradient(loss_fun, data, weights)  
    weights += - step_size * weights_grad # perform parameter update
```



Stochastic Gradient Descent

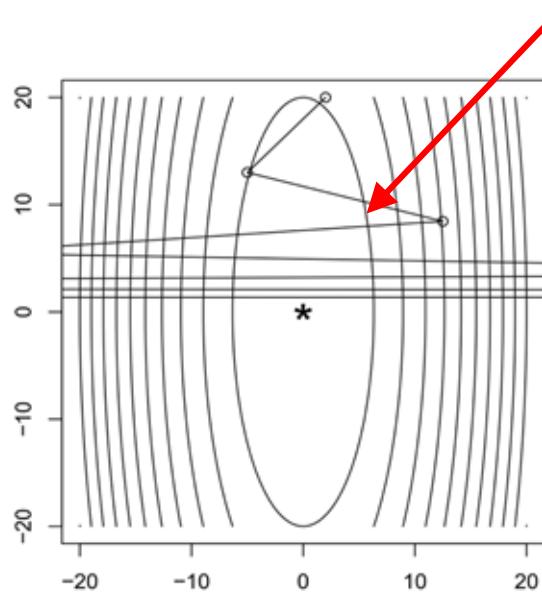
- Goal:
 - Optimize parameters to minimize loss
- Step along the direction of steepest descent (negative gradient)

$$\boxed{\theta_t} = \boxed{\theta_{t-1}} - \boxed{\alpha} \boxed{\nabla_{\theta_{t-1}} f(\theta_{t-1})}$$

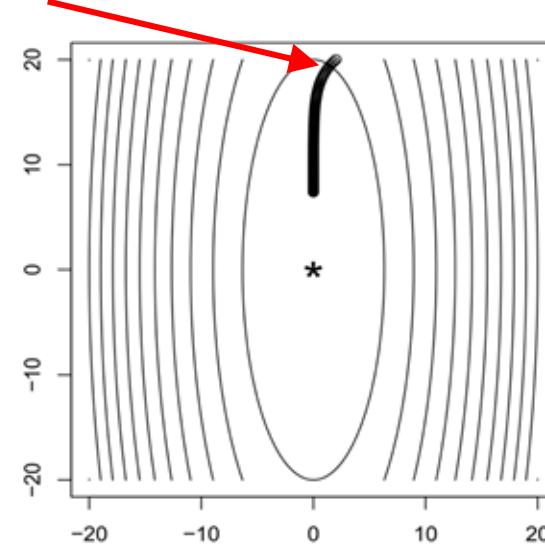


What does SGD do?

- Gradient over **entire dataset** is **impractical**
- Better to **take quick, noisy steps**
- Go fast in **along the consistent direction**
- Estimate **gradient over a mini-batch of samples**
 - Increasing batch **size on GPU** is **essentially free up to a point**



Too big learning rate



Too small learning rate

SGD with Momentum

$$\boxed{\theta_t} = \boxed{\theta_{t-1}} - \boxed{\alpha} \boxed{\nabla_{\theta_{t-1}} f(\theta_{t-1})}$$

$$\boxed{v_t} = \mu \boxed{v_{t-1}} - \boxed{\alpha} \boxed{\nabla f(\theta_{t-1})}$$

$$\boxed{\theta_t} = \boxed{\theta_{t-1}} + \boxed{v_t}$$

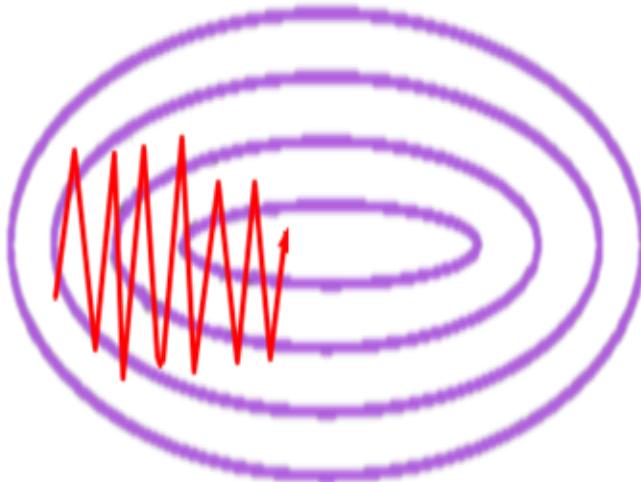
$$v_t = \boxed{\mu} v_{t-1} - \alpha \nabla f(\theta_{t-1})$$

$$\theta_t = \theta_{t-1} + v_t$$

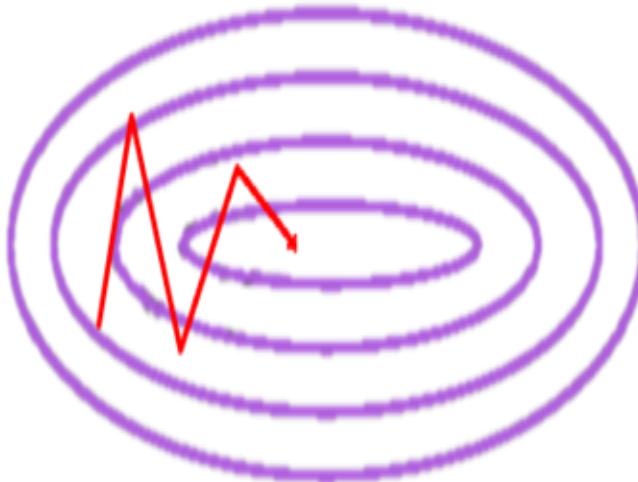
$$\theta_t = \theta_{t-1} + (\boxed{0} \cdot v_{t-1} - \alpha \nabla f(\theta_{t-1}))$$

Same as vanilla gradient descent

Impact of SGD with momentum



SGD without Momentum



SGD with Momentum

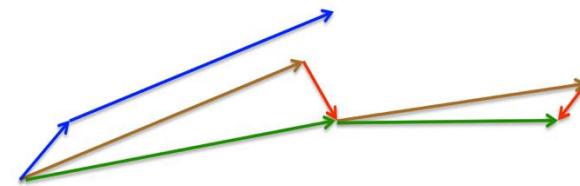
- **Cancels out oscillation**
- **Gathers speed** in direction that matters
- We can adapt our **updates to the slope of our error function and speed up SGD in turn**

Nesterov's Accelerated Gradient (NAG)

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$

$$\theta = \theta - v_t$$

γ is usually set to 0.9 or a similar value



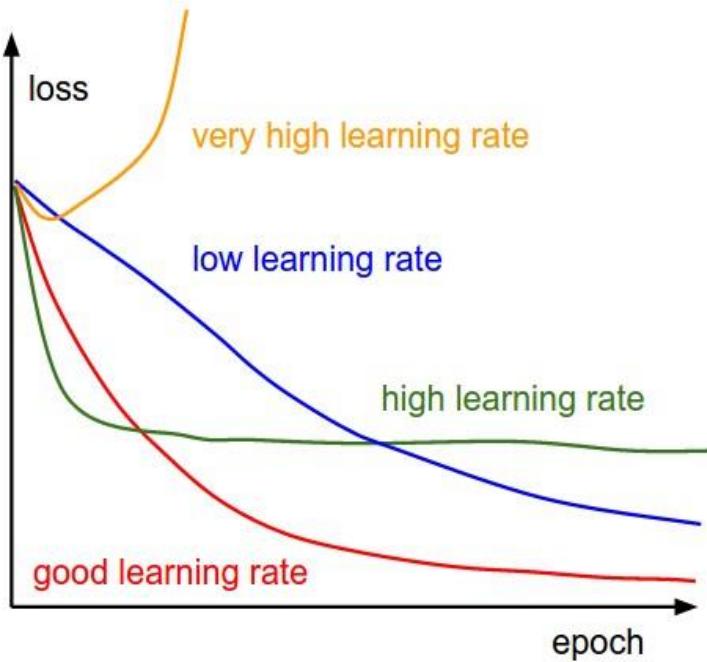
Nesterov update (Source: [G. Hinton's lecture 6c](#))

- First computes the **current gradient (small blue vector in figure)** and then takes a **big jump in the direction of the updated accumulated gradient (big blue vector)**
- NAG first makes a big jump in the **direction of the previous accumulated** gradient (brown vector), measures the gradient and then makes a correction (red vector), which results in the **complete NAG update (green vector)**.
- Advantages:
 - Prevents from going too fast and results in increased responsiveness
 - Significantly increased the performance of RNNs

Challenges : per parameter learning rate

- Gradients of **different layers have different magnitudes**
- Different neuron units **have different firing rates**
- So, what we want :
 - We would also like to adapt our updates to each individual parameter to **perform larger or smaller updates depending on their importance** (different learning rates for different parameters)

How to pick the learning rate?



- **Too big = diverge**, too small = slow convergence
- No “one learning rate to rule them all”
- Start from a high value and **keep cutting by half if model diverges**
- Learning rate schedule:
 - decay learning rate over time

Adagrad

- **Adagrad** adapts the learning rate to the parameters, performing
 - Smaller updates (low learning rates) for parameters associated with frequently occurring features, and
 - Larger updates (high learning rates) for parameters associated with infrequent features.
- But how?
 - Gradient update depends on **history of magnitude of gradients**
 - Parameters with **small or sparse updates have larger learning rates**
 - **Square root** important for good performance
 - More **tolerance for learning rate**

Adagrad

$$\theta_t = \theta_{t-1} - \alpha \nabla_{\theta_{t-1}} f(\theta_{t-1})$$

$$\begin{aligned} g_t &= \sum_{\tau=1}^{t-1} (\nabla f(\theta_\tau))^2 \\ &= g_{t-1} + (\nabla f(\theta_{t-1}))^2 \end{aligned}$$

$$\theta_t = \theta_{t-1} - \frac{\alpha}{\sqrt{g_t + \epsilon}} \odot \nabla f(\theta_{t-1})$$

$$\begin{aligned} \boxed{g_t} &= \sum_{\tau=1}^{t-1} (\nabla f(\theta_\tau))^2 \\ &= g_{t-1} + (\nabla f(\theta_{t-1}))^2 \end{aligned}$$

What happens as t increases?

$$\theta_t = \theta_{t-1} - \frac{\alpha}{\sqrt{\boxed{g_t} + \epsilon}} \odot \nabla f(\theta_{t-1})$$

Adagrad

- Advantages :
 - **Eliminates** the need to **manually tune the learning rate**.
 - Most implementations use **a default value of 0.01**
 - **Ensure faster and better training process**
- **Dis-advantages:**
 - **Maintain entire history of gradients**
 - As a results, the **sum of magnitude of gradients always increasing** during training process.
 - **Learning rate to shrink and eventually become infinitesimally small**, at which point the algorithm is **no longer able to acquire additional knowledge** (Forces learning rate to 0 over time)
 - **Hard to compensate** for in advance

Solution: don't maintain all history

- **Monotonically increasing** because we hold all the history instead, **forget gradients far in the past**
- In practice, **downweight previous gradients exponentially**
- The following algorithms aim to resolve this flaw.
 - **AdaDelta, and**
 - **RMSProp**

Adadelta

- **Adadelta** is an extension of Adagrad that seeks to **reduce its aggressive, monotonically decreasing learning rate**
- Instead of accumulating all past squared gradients, Adadelta **restricts the window of accumulated past gradients to some fixed size w**
- The running average $E[g^2]_t$ at time step t then depends (as a fraction γ similarly to the Momentum term) **only on the previous average and the current gradient**

$$E[g^2]_t = \gamma E[g^2]_{t-1} + (1 - \gamma) g_t^2.$$

γ to a similar value as the momentum term, around 0.9.

Adadelta

- What and where to replace:

$$\Delta\theta_t = -\frac{\eta}{\sqrt{G_t + \epsilon}} \odot g_t. \quad \Delta\theta_t = -\frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t.$$

- As the denominator is just the **root mean squared (RMS)** error criterion of the gradient, we can replace it with the criterion short-hand:

$$\Delta\theta_t = -\frac{\eta}{RMS[g]_t} g_t.$$

RMSProp

- RMSprop is an unpublished, **adaptive learning rate method proposed by Geoff Hinton**
- RMSprop and Adadelta have both been developed independently around the same time stemming from the need to **resolve Adagrad's radically diminishing learning rates**.
- RMSprop **only cares about recent gradients** (otherwise like Adagrad)
 - Good property because **optimization landscape changes**
- Standard **gamma (γ) is 0.9**

$$\begin{aligned} g_t &= \sum_{\tau=1}^{t-1} (\nabla f(\theta_\tau))^2 \\ &= g_{t-1} + (\nabla f(\theta_{t-1}))^2 \end{aligned}$$

$$\begin{aligned} g_t &= (1 - \boxed{\gamma}) \sum_{\tau=1}^{t-1} \boxed{\gamma^{t-1-\tau}} (\nabla f(\theta_\tau))^2 \\ &= \boxed{\gamma} g_{t-1} + \boxed{(1 - \gamma)} \nabla(f(\theta_{t-1}))^2 \end{aligned}$$

$$\begin{aligned}
g_t &= (1 - \gamma) \sum_{\tau=1}^{t-1} \gamma^{t-1-\tau} (\nabla f(\theta_\tau))^2 \\
&= \gamma g_{t-1} + (1 - \gamma) \nabla(f(\theta_{t-1}))^2
\end{aligned}$$

$$\theta_t = \theta_{t-1} - \frac{\alpha}{\sqrt{g_t + \epsilon}} \odot \nabla f(\theta_{t-1})$$

- RMSprop also divides the learning rate by an exponentially decaying average of squared gradients.
- **Hinton suggests γ to be set to 0.9**, while a good default value for the learning rate η is 0.001.

Adam optimizer

- Adam is derived from **Adaptive Moment Estimation (Adam)** for each parameter
- Adam also keeps an exponentially decaying average of past gradients m_t , similar to momentum
- Essentially, **combine RMSProp + Momentum**
- Includes bias correction term from **initializing m and v to 0**

$$\boxed{\mu}v_{t-1} - \boxed{\alpha}\nabla f(\theta_{t-1}) \quad \text{Momentum}$$

$$\boxed{\gamma}g_{t-1} + \boxed{(1 - \gamma)}\nabla(f(\theta_{t-1}))^2 \quad \text{RMSProp}$$

$$m_t = \boxed{\beta_1}m_{t-1} + \boxed{(1 - \beta_1)}\nabla f(\theta_{t-1})$$

$$v_t = \boxed{\beta_2}v_{t-1} + \boxed{(1 - \beta_2)}(\nabla f(\theta_{t-1}))^2$$

Ref. : The relationship between second moment $E[X^2]$ and the variance $Var(X)$:

$$\begin{aligned} Var(X) &= E[X^2] - (E[X])^2 \\ E[X^2] &= Var(X) + (E[X])^2 \end{aligned}$$

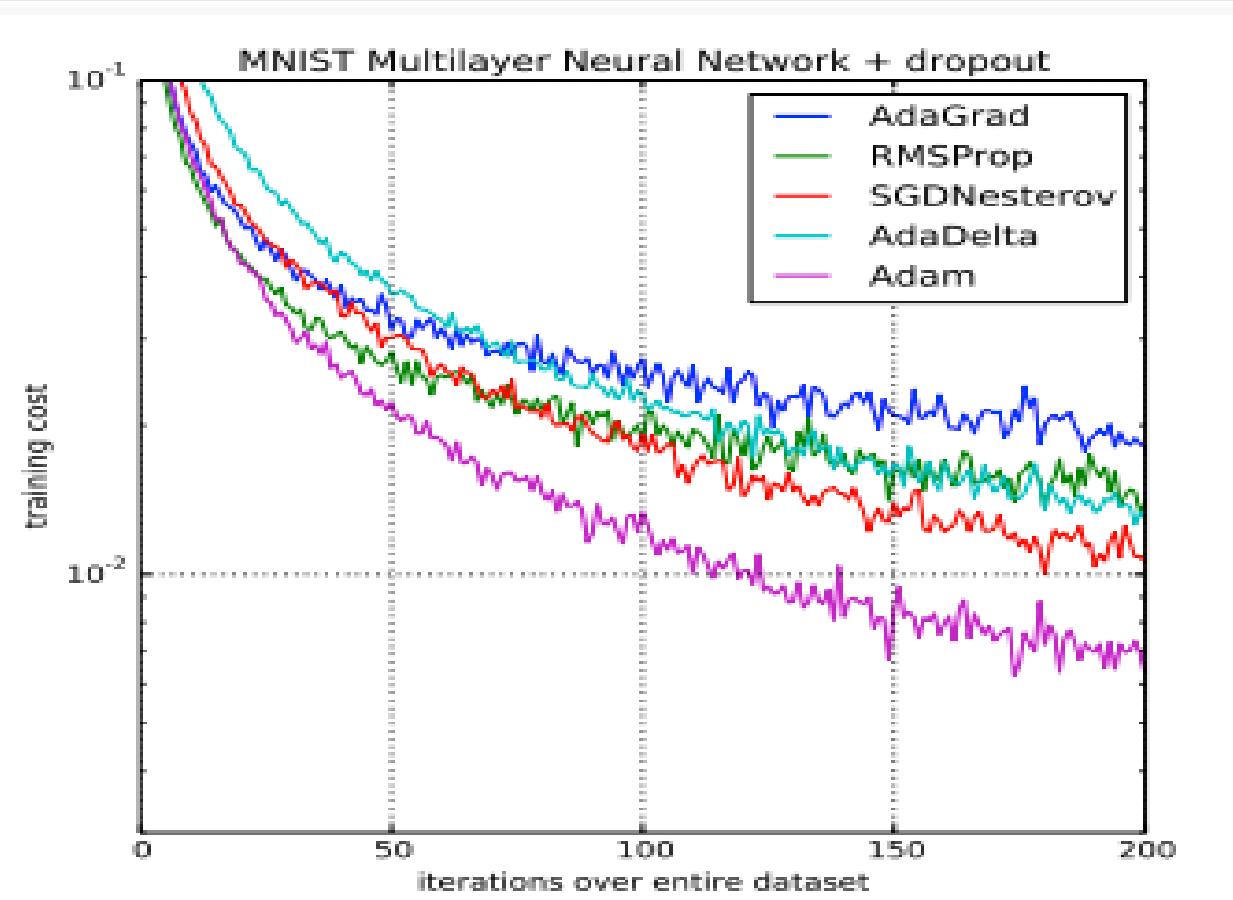
$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla f(\theta_{t-1})$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) (\nabla f(\theta_{t-1}))^2$$

$$\theta_t = \theta_{t-1} - \frac{\alpha}{\sqrt{v_t + \epsilon}} \odot m_t$$

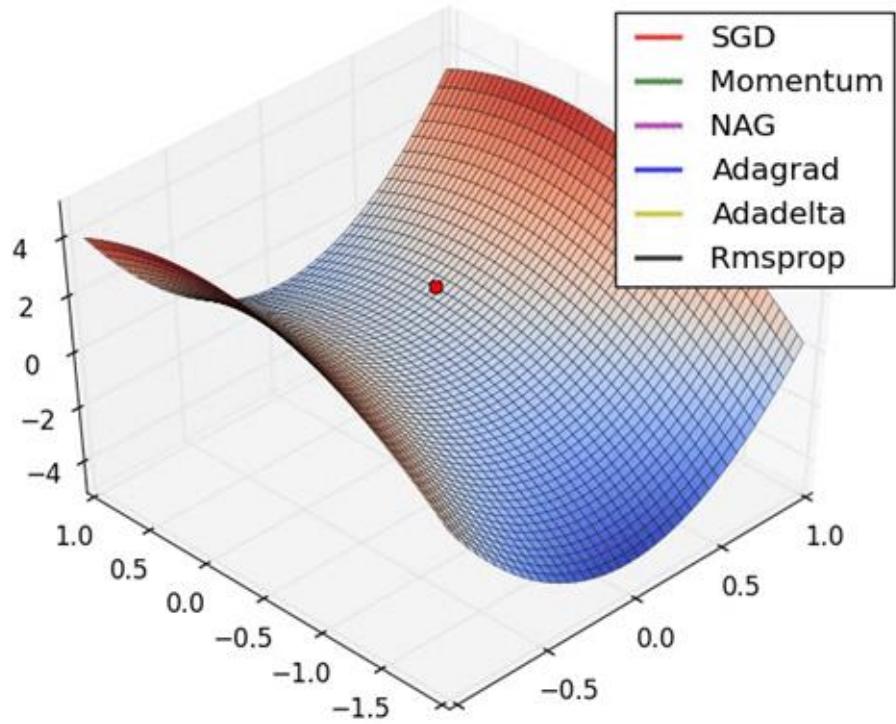
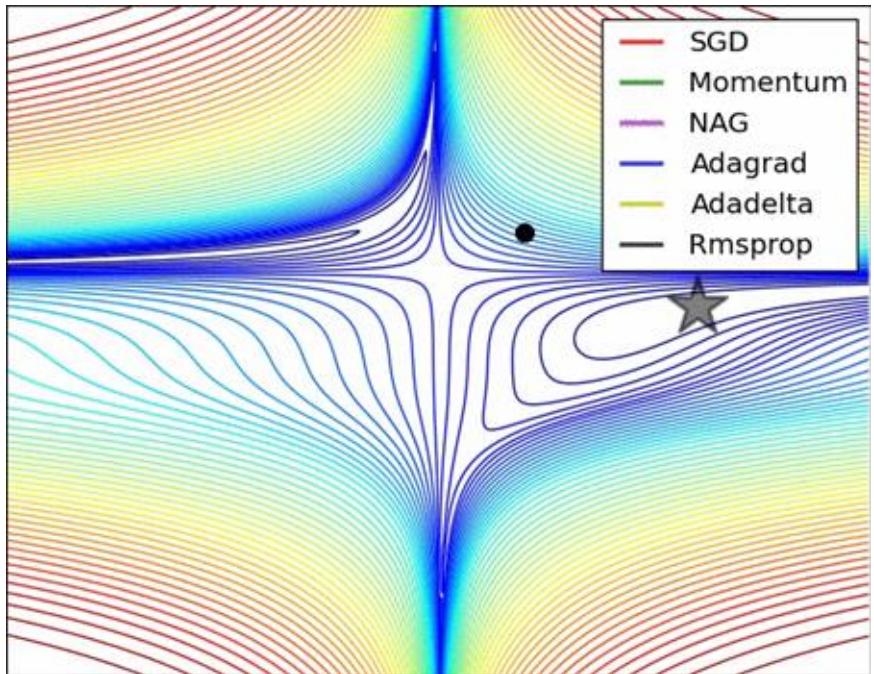
As m_t and v_t are initialized as vectors of 0's, the authors of Adam observed that **they are biased towards zero**, especially during the initial time steps, and when the decay rates are small (i.e. β_1 and β_2 are close to 1)

Comparison



Comparison of Adam to Other Optimization Algorithms Training
a Multilayer Perceptron

Taken from Adam: A Method for Stochastic Optimization, 2015.



Alec Radford

Why AdamW matters

- The idea behind **L2 regularization or weight decay** is that networks with smaller weights (all other things being equal) are observed
- Help to overfit less and generalize better**

Algorithm 2 Adam with L₂ regularization and

Adam with weight decay (AdamW)

- 1: **given** $\alpha = 0.001, \beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}, w \in \mathbb{R}$
- 2: **initialize** time step $t \leftarrow 0$, parameter vector $\mathbf{x}_{t=0} \in \mathbb{R}^n$, first moment vector $\mathbf{m}_{t=0} \leftarrow \mathbf{0}$, second moment vector $\mathbf{v}_{t=0} \leftarrow \mathbf{0}$, schedule multiplier $\eta_{t=0} \in \mathbb{R}$
- 3: **repeat**
- 4: $t \leftarrow t + 1$
- 5: $\nabla f_t(\mathbf{x}_{t-1}) \leftarrow \text{SelectBatch}(\mathbf{x}_{t-1})$ ▷ select batch and return the corresponding gradient
- 6: $\mathbf{g}_t \leftarrow \nabla f_t(\mathbf{x}_{t-1}) + w\mathbf{x}_{t-1}$
- 7: $\mathbf{m}_t \leftarrow \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t$ ▷ here and below all operations are element-wise
- 8: $\mathbf{v}_t \leftarrow \beta_2 \mathbf{v}_{t-1} + (1 - \beta_2) \mathbf{g}_t^2$
- 9: $\hat{\mathbf{m}}_t \leftarrow \mathbf{m}_t / (1 - \beta_1^t)$ ▷ β_1 is taken to the power of t
- 10: $\hat{\mathbf{v}}_t \leftarrow \mathbf{v}_t / (1 - \beta_2^t)$ ▷ β_2 is taken to the power of t
- 11: $\eta_t \leftarrow \text{SetScheduleMultiplier}(t)$ ▷ can be fixed, decay, or also be used for warm restarts
- 12: $\mathbf{x}_t \leftarrow \mathbf{x}_{t-1} - \eta_t \left(\alpha \hat{\mathbf{m}}_t / (\sqrt{\hat{\mathbf{v}}_t} + \epsilon) + w\mathbf{x}_{t-1} \right)$
- 13: **until** stopping criterion is met
- 14: **return** optimized parameters \mathbf{x}_t

Other recent optimizers

- **AMSGrad**
- QHAdam
- AggMo
- **EVE**
- What and how to use it?
 - Adaptive optimizers like Adam have become **a default choice** for training neural networks
 - Researchers often prefer stochastic gradient descent (**SGD**) with **momentum** because models **trained with Adam have been observed to not generalize as well**
 - **SGD + momentum and Adam** are good first steps
 - Just use **Adam with default parameters**
 - **Learning rate decay** always good

1.Loshchilov, I., & Hutter, F. (2019). Decoupled Weight Decay Regularization. In Proceedings of ICLR 2019. [↗](#)

2.Ma, J., & Yarats, D. (2019). Quasi-hyperbolic momentum and Adam for deep learning. In Proceedings of ICLR 2019. [↗](#)

3.Lucas, J., Sun, S., Zemel, R., & Grosse, R. (2019). Aggregated Momentum: Stability Through Passive Damping. In Proceedings of ICLR 2019. [↗](#)

4.Niu, F., Recht, B., Christopher, R., & Wright, S.J. (2011). Hogwild! : A Lock-Free Approach to Parallelizing Stochastic Gradient Descent, 1–22.

Recent development: why new optimizers?

- Large deep models (LLMs, multimodal learning) suffer from:
 - Poor curvature awareness ($\text{Adam} \neq \text{true second order}$)
 - Instability at scale
 - Overconfidence and sharp minima (important for clinical AI)
- **Sophia and Muon** are **next-generation optimizers** designed to approximate **second-order information efficiently** while remaining scalable.

Sophia (Second-order Clipped Stochastic Optimization) Optimizer

Classic optimizers:

| Optimizer | Uses gradient | Uses curvature |
|-----------|---------------|----------------|
| SGD | ✓ | ✗ |
| Adam | ✓ (adaptive) | ✗ |
| Newton | ✓ | ✓ (Hessian) |

Problem:

True Newton updates require computing or inverting the Hessian H , which is infeasible for deep networks.

Sophia approximates **diagonal Hessian information** using stochastic estimates, yielding:

Second-order behavior at Adam-like cost

Mathematical Foundation

Objective

$$\min_{\theta} \quad \mathcal{L}(\theta)$$

Taylor expansion

$$\mathcal{L}(\theta + \Delta) \approx \mathcal{L}(\theta) + g^T \Delta + \frac{1}{2} \Delta^T H \Delta$$

Newton step:

$$\Delta = -H^{-1}g$$

1

Sophia's Key Approximation

Sophia replaces H with a **diagonal estimate**:

$$H \approx \text{diag}(h_1, \dots, h_d)$$

Each diagonal term is estimated via **gradient outer products**:

$$h_t \approx \mathbb{E}[g_t^2]$$

This is *similar to Adam*, but **interpreted as curvature**, not variance.

2

Sophia Update Rule

Let:

- $g_t = \nabla_{\theta} \mathcal{L}$
- h_t = running estimate of Hessian diagonal
- m_t = momentum

Hessian estimate

$$h_t = \beta h_{t-1} + (1 - \beta) g_t^2$$

Parameter update

$$\theta_{t+1} = \theta_t - \eta \cdot \frac{m_t}{\max(h_t, \epsilon)}$$

Why Sophia is powerful

Key difference from Adam

Adam

g^2 =variance

Scale learning rate

Sophia

g^2 =curvature

Scale Newton step

- Approximates **diagonal Newton method**
- Naturally **prevents sharp minima**
- Improves **generalization**
- Works extremely well for:
 - LLMs
 - Vision Transformers
 - Multimodal Learning

Muon Optimizer

Motivation : Sophia handles **curvature**, but large models also suffer from:

- Gradient misalignment across layers
- Poor signal propagation
- Feature collapse

Muon introduces geometry-aware updates preserve **directional structure** in parameter space.

Core Idea : treats optimization as **vector transport in parameter space**, not just **scalar learning rates**. It enforces:

- Orthogonality
- Directional consistency
- Controlled norm growth

Mathematical Intuition and update rule

1

Let gradient matrix $G \in \mathbb{R}^{d \times k}$ (e.g., attention layers).

Muon applies **orthogonal projection**:

$$\tilde{G} = \text{Proj}_{\mathcal{O}}(G)$$

This ensures:

$$\tilde{G}^\top \tilde{G} \approx I$$

Meaning updates preserve **feature diversity**.

2

Step 1: Normalize gradient direction

$$\hat{g} = \frac{g}{\| g \|}$$

Step 2: Apply momentum and geometry correction

$$\theta_{t+1} = \theta_t - \eta \cdot \hat{g}$$

In practice, Muon uses **block-wise normalization** for:

- Attention heads
- Convolution kernels
- Transformer layers

Muon is an optimizer for 2D parameters of neural network hidden layers

Algorithm 2 Muon

Require: Learning rate η , momentum μ

- 1: Initialize $B_0 \leftarrow 0$
 - 2: **for** $t = 1, \dots$ **do**
 - 3: Compute gradient $G_t \leftarrow \nabla_{\theta} \mathcal{L}_t(\theta_{t-1})$
 - 4: $B_t \leftarrow \mu B_{t-1} + G_t$
 - 5: $O_t \leftarrow \text{NewtonSchulz5}(B_t)$
 - 6: Update parameters $\theta_t \leftarrow \theta_{t-1} - \eta O_t$
 - 7: **end for**
 - 8: **return** θ_t
-

```
# Pytorch code
def newtonschulz5(G, steps=5, eps=1e-7):
    assert G.ndim == 2
    a, b, c = (3.4445, -4.7750, 2.0315)
    X = G.bfloat16()
    X /= (X.norm() + eps)
    if G.size(0) > G.size(1):
        X = X.T
    for _ in range(steps): A = X @ X.T
        B = b * A + c * A @ A
        X = a * X + B @ X
    if G.size(0) > G.size(1):
        X = X.T
    return X
```

Sophia vs Muon (Comparison)

| Aspect | Sophia | Muon |
|--------------|-----------------------------|------------------------|
| Focus | Curvature | Geometry |
| Second-order | Approximate Hessian | Directional control |
| Main benefit | Faster + stable convergence | Better representations |
| Best for | LLMs, classification | Transformers, fusion |

Sophia + Muon hybrid:

- Sophia for **encoder weights**
- Muon for **fusion & attention layers**

Benefits : Newton for stability and Riemannian optimization for geometry

Summary

- Efficient training of DNN models
 - Loss functions:
 - Cross-entropy Loss
 - Zero-one loss
 - Binary Cross Entropy Loss and
 - Categorical Cross Entropy Loss
 - Focal Loss
- Optimization methods
 - SGD + momentum
 - RMSprop
 - Adam
 - AdamW
 - Sophia and muon many more ...
- What's next?
 - Deep Neural Network architectures