

COMP/EECE 7/8740 Neural Networks

Topics:

- Receptive field for deep CNN (DCNN)
- CNN architectures or models
 - Classification
 - Segmentation (regression)
 - Detection (classification + regression)
- Design principle of DCNN architectures


Md Zahangir Alom
Department of Computer Science
University of Memphis, TN

Importance of Studying DCNN Architectures

- DCNN architectures capture the **key design principles**—
 - Depth, residual connections, attention, and efficient scaling
- Architectures are the **major breakthroughs** in
 - Accuracy, efficiency, and generalization across tasks
- Foundation of modern applications and research,
 - Enabling practitioners to **design better models**


Receptive field in CNNs

- The **receptive field** is defined as the region in the input space that a particular CNN's feature is looking at (i.e. be affected by).
- **Not all pixels in a receptive field is equally important** to its corresponding CNN's feature
- Closer a pixel to the **center of the RF, the more it contributes to the calculation** of the output feature (focus exponentially more to the middle of that region).



Convolution and receptive field


Receptive field in CNNs



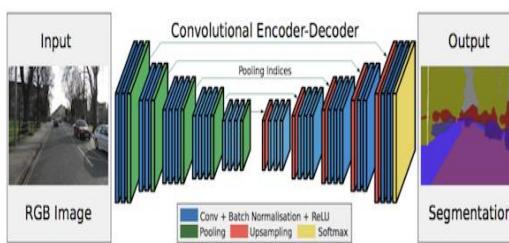
- By applying a convolution C with kernel size $k = 3 \times 3$, padding size $p = 1 \times 1$, stride $s = 2 \times 2$ on an input map 5×5 , we will get an output feature map 3×3 (green map).
- Applying the same convolution on top of the 3×3 feature map, we will get a 2×2 feature map (orange map).

Deep CNN Models/Architectures


▪ Classification



▪ Segmentation



▪ Detection



Models for classification:

- AlexNet
- VGG Net
- GoogleNet
- ResNet
- Inception-ResNet
- DenseNet / DCRN
- FractalNet
- CapsuleNet and
- IRRCNN
-
- Transformer Networks

Models for Segmentation:

- FCN
- SegNet
- Dilaited Convolution
- RefineNet
- Pyramid Scene Parsing (PSP):PSPNet
- DeepLab
- U-Net
- R2U-Net
- NABLA-N Net
- Segment Anything Model(SAM)
-

Models for Detection:

- Region based CNN (RCNN)
- Fast RCNN
- Faster RCNN
- Mask RCNN
- You Only Look Once (YOLO)
- Single Short Multibox Detector (SSD)
- UD-Net
-


CNN models

- **LeNet** : Yann LeCun in 1998
- **AlexNet** : ILSVR winner in 2012
- **ZFNet** : Matthew Zeiler and Rob Fergue won the ILSVRC 2013, Refinement of AlexNet
- **VGGNET**: Visual Geometry Group (VGG) from Oxford University runner up of ILSVRC in 2014.
- **Network in Network (NiN)**: from NUS in 2014

CNN models


- **GoogLeNet(2014)**: Szegedy from the Google who was the winner of ILSVRC in 2014.
- **ResNet (2015)** : from Microsoft won the ILSVRC in 2015.
- **Inception-Residual Network by C. Szegedy** in 2016
- **DenseNet (Dec. 2016)** : from Cornel University by Gao Huang and others (CVPR-2017 best paper award) from Cornell University
- **FractalNet (2016)**: Ultra-Deep Neural Networks without Residuals from University of Chicago.
- **PolyNet**
- **Res2Net in 2019**
- **ConvNext**
- **Transformer (such as ViT) and**
- **JEPA**

LeNet-5



- Average pooling
- Sigmoid or tanh nonlinearity
- Fully connected layers at the end
- Trained on MNIST digit dataset with 60K training examples


AlexNet: ILSVRC 2012 winner



- Similar Framework to LeNet but:
 - Max pooling, ReLU nonlinearity
 - More data and bigger model (7 hidden layers, 650K unit, 61M params)
 - GPU implementation (50x speed up over CPU)
 - Trained on two GPUs for a week
 - Dropout regularization
 - **Local Response Normalization (LRN)**

Local Response Normalization (LRN)

- LRN layer implements the lateral inhibition and **objective to amplify the excited neuron** while dampening the surrounding neurons.
- Two approaches for LRN:
 - Consider same channel or feature map and 2D neighborhood of dimension $N \times N$, where N is the size of the normalization window. **Normalize the window using the values in this neighborhood.**
 - **Normalizing across channels or feature maps**, you will consider a neighborhood along the third dimension but at a single location.




First approach



Second approach

AlexNet: ILSVRC 2012 winner


- The size of input sample is 224x224x3,
- Filter/ Kernel/ weight size 11,
- Stride 4 and the
- Output of the first convolution layer is 55x55x96.



(Krizhevsky NIPS 2014)


Summary on AlexNet :2012

- The **first deep learning model** shown to be effective on large scale computer vision task.
- The first time a **very large scale** deep model is adopted.
- **GPU is shown to be every effective** on this large scale deep learning model.



How it's implemented for ImageNet

- ImageNet are divided into two groups of 500 classes, A and B
- Two 8-layer AlexNets, base A and base B, are trained on the two groups respectively



Clarifai/ ZFNet : ILSVRC 2013 winner

- Refinement of AlexNet





Figure 3. Architecture of our 8 layer convnet model. A 224 by 224 crop of an image (with 3 color planes) is presented as the input. This is convolved with 96 different 1st layer filters (red), each of size 7 by 7, using a stride of 2 in both x and y. The resulting feature maps are then: (i) passed through a rectified linear function (not shown), (ii) pooled (max within 3x3 regions, using stride 2) and (iii) contrast normalized across feature maps to give 96 different 55 by 55 element feature maps. Similar operations are repeated in layers 2,3,4,5. The last two layers are fully connected, taking features from the top convolutional layer as input in vector form ($6 \cdot 6 \cdot 256 = 9216$ dimensions). The final layer is a C -way softmax function, C being the number of classes. All filters and feature maps are square in shape.

Clarifai/ZFNet: ILSVRC 2013 winner

- Max-pooling layers follow first, second, and fifth convolutional layers
- 11*11 to 7*7, stride 4 to 2 in 1st layer (increasing resolution of feature maps)
- Other settings are the same as AlexNet
- Reduce the error by 2%.

M. Zeiler and R. Fergus, [Visualizing and Understanding Convolutional Networks](#),
ECCV 2014 (Best Paper Award winner)

VGGNet: ILSVRC 2014 2nd place



K. Simonyan and A. Zisserman, [Very Deep Convolutional Networks for Large-Scale Image Recognition](#), ICLR 2015


VGGNet: ILSVRC 2014 2nd place

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224×224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: Number of parameters (in millions).

Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

- Sequence of deeper networks trained progressively
- Large receptive fields replaced by successive layers of 3x3 convolutions (with ReLU in between)



- One 7x7 conv layer with C size of feature maps needs $49C^2$ weights, three 3x3 conv layers need only $27C^2$ weights
- Experimented with 1x1 convolutions

VGGNet: ILSVRC 2014 2nd place

- INPUT: [224x224x3] memory: $224 \times 224 \times 3 = 150K$ params: 0
- CONV3-64: [224x224x64] memory: $224 \times 224 \times 64 = 3.2M$ params: $(3 \times 3 \times 3) \times 64 = 1,728$
- CONV3-64: [224x224x64] memory: $224 \times 224 \times 64 = 3.2M$ params: $(3 \times 3 \times 64) \times 64 = 36,864$
- POOL2: [112x112x64] memory: $112 \times 112 \times 64 = 800K$ params: 0
- CONV3-128: [112x112x128] memory: $112 \times 112 \times 128 = 1.6M$ params: $(3 \times 3 \times 64) \times 128 = 73,728$
- CONV3-128: [112x112x128] memory: $112 \times 112 \times 128 = 1.6M$ params: $(3 \times 3 \times 128) \times 128 = 147,456$
- POOL2: [56x56x128] memory: $56 \times 56 \times 128 = 400K$ params: 0
- CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800K$ params: $(3 \times 3 \times 128) \times 256 = 294,912$
- CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800K$ params: $(3 \times 3 \times 256) \times 256 = 589,824$
- CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800K$ params: $(3 \times 3 \times 256) \times 256 = 589,824$
- POOL2: [28x28x256] memory: $28 \times 28 \times 256 = 200K$ params: 0
- CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400K$ params: $(3 \times 3 \times 256) \times 512 = 1,179,648$
- CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400K$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
- CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400K$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
- POOL2: [14x14x512] memory: $14 \times 14 \times 512 = 100K$ params: 0
- CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100K$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
- CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100K$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
- CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100K$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
- POOL2: [7x7x512] memory: $7 \times 7 \times 512 = 25K$ params: 0
- FC: [1x1x4096] memory: 4096 params: $7 \times 7 \times 512 \times 4096 = 102,760,448$
- FC: [1x1x4096] memory: 4096 params: $4096 \times 4096 = 16,777,216$
- FC: [1x1x1000] memory: 1000 params: $4096 \times 1000 = 4,096,000$

	D	E
1	16 weight layers	19 weight layers
age 0		
1	conv3-64	conv3-64
1	conv3-64	conv3-64
2		
3	conv3-128	conv3-128
3	conv3-128	conv3-128
4		
5	conv3-256	conv3-256
5	conv3-256	conv3-256
6	conv3-256	conv3-256
6		conv3-256
7		
2	conv3-512	conv3-512
2	conv3-512	conv3-512
2	conv3-512	conv3-512
2		conv3-512
2		
2	conv3-512	conv3-512
2	conv3-512	conv3-512
2	conv3-512	conv3-512
2		

VGGNet: ILSVRC 2014 2nd place

- INPUT: [224x224x3] **memory: 224*224*3=150K** params: 0
- CONV3-64: [224x224x64] **memory: 224*224*64=3.2M** params: $(3*3*3)*64 = 1,728$
- CONV3-64: [224x224x64] **memory: 224*224*64=3.2M** params: $(3*3*64)*64 = 36,864$
- POOL2: [112x112x64] **memory: 112*112*64=800K** params: 0
- CONV3-128: [112x112x128] **memory: 112*112*128=1.6M** params: $(3*3*64)*128 = 73,728$
- CONV3-128: [112x112x128] **memory: 112*112*128=1.6M** params: $(3*3*128)*128 = 147,456$
- POOL2: [56x56x128] **memory: 56*56*128=400K** params: 0
- CONV3-256: [56x56x256] **memory: 56*56*256=800K** params: $(3*3*128)*256 = 294,912$
- CONV3-256: [56x56x256] **memory: 56*56*256=800K** params: $(3*3*256)*256 = 589,824$
- CONV3-256: [56x56x256] **memory: 56*56*256=800K** params: $(3*3*256)*256 = 589,824$
- POOL2: [28x28x256] **memory: 28*28*256=200K** params: 0
- CONV3-512: [28x28x512] **memory: 28*28*512=400K** params: $(3*3*256)*512 = 1,179,648$
- CONV3-512: [28x28x512] **memory: 28*28*512=400K** params: $(3*3*512)*512 = 2,359,296$
- CONV3-512: [28x28x512] **memory: 28*28*512=400K** params: $(3*3*512)*512 = 2,359,296$
- POOL2: [14x14x512] **memory: 14*14*512=100K** params: 0
- CONV3-512: [14x14x512] **memory: 14*14*512=100K** params: $(3*3*512)*512 = 2,359,296$
- CONV3-512: [14x14x512] **memory: 14*14*512=100K** params: $(3*3*512)*512 = 2,359,296$
- CONV3-512: [14x14x512] **memory: 14*14*512=100K** params: $(3*3*512)*512 = 2,359,296$
- POOL2: [7x7x512] **memory: 7*7*512=25K** params: 0
- FC: [1x1x4096] **memory: 4096** params: $7*7*512*4096 = 102,760,448$
- FC: [1x1x4096] **memory: 4096** params: $4096*4096 = 16,777,216$
- FC: [1x1x1000] **memory: 1000** params: $4096*1000 = 4,096,000$

TOTAL memory: 24M * 4 bytes \sim 93MB / image (only forward! \sim *2 for bwd)

TOTAL params: 138M parameters

VGGNet: ILSVRC 2014 2nd place

- INPUT: [224x224x3] memory: $224 \times 224 \times 3 = 150\text{K}$ params: 0
- CONV3-64: [224x224x64] memory: $224 \times 224 \times 64 = 3.2\text{M}$ params: $(3 \times 3 \times 3) \times 64 = 1,728$
- CONV3-64: [224x224x64] memory: $224 \times 224 \times 64 = 3.2\text{M}$ params: $(3 \times 3 \times 64) \times 64 = 36,864$
- POOL2: [112x112x64] memory: $112 \times 112 \times 64 = 800\text{K}$ params: 0
- CONV3-128: [112x112x128] memory: $112 \times 112 \times 128 = 1.6\text{M}$ params: $(3 \times 3 \times 64) \times 128 = 73,728$
- CONV3-128: [112x112x128] memory: $112 \times 112 \times 128 = 1.6\text{M}$ params: $(3 \times 3 \times 128) \times 128 = 147,456$
- POOL2: [56x56x128] memory: $56 \times 56 \times 128 = 400\text{K}$ params: 0
- CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800\text{K}$ params: $(3 \times 3 \times 128) \times 256 = 294,912$
- CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800\text{K}$ params: $(3 \times 3 \times 256) \times 256 = 589,824$
- CONV3-256: [56x56x256] memory: $56 \times 56 \times 256 = 800\text{K}$ params: $(3 \times 3 \times 256) \times 256 = 589,824$
- POOL2: [28x28x256] memory: $28 \times 28 \times 256 = 200\text{K}$ params: 0
- CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400\text{K}$ params: $(3 \times 3 \times 256) \times 512 = 1,179,648$
- CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400\text{K}$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
- CONV3-512: [28x28x512] memory: $28 \times 28 \times 512 = 400\text{K}$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
- POOL2: [14x14x512] memory: $14 \times 14 \times 512 = 100\text{K}$ params: 0
- CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100\text{K}$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
- CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100\text{K}$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
- CONV3-512: [14x14x512] memory: $14 \times 14 \times 512 = 100\text{K}$ params: $(3 \times 3 \times 512) \times 512 = 2,359,296$
- POOL2: [7x7x512] memory: $7 \times 7 \times 512 = 25\text{K}$ params: 0
- FC: [1x1x4096] memory: 4096 params: $7 \times 7 \times 512 \times 4096 = 102,760,448$
- FC: [1x1x4096] memory: 4096 params: $4096 \times 4096 = 16,777,216$
- FC: [1x1x1000] memory: 1000 params: $4096 \times 1000 = 4,096,000$


Most memory is in early CONV

Most params are in late FC


TOTAL memory: $24\text{M} * 4 \text{ bytes} \approx 93\text{MB} / \text{image}$ (only forward! ≈ 2 for bwd)

TOTAL params: 138M parameters


Network in network:NUS-2013





(a) Linear convolution layer



(b) Mlpconv layer




1×1 convolutions



1×1 conv layer

Advantages on NiN

- Remove the two fully connected layers (fc6, fc7) of the AlexNet but add NIN into the AlexNet.



	Parameter Number	Performance	Time to train (GTX Titan)
AlexNet	60 Million (230 Megabytes)	40.7% (Top 1)	8 days
NIN	7.5 Million (29 Megabytes)	39.2% (Top 1)	4 days

Advantages of using 1x1 Conv.

- **Reduce the number of computational** parameters
- **Add more non-linearity** on multiple levels of feature representation without dimensionality reduction
- **Mapping on any number of feature maps** (higher to lower and lower to higher dimension)



WE NEED TO GO

DEEPER

“Understanding” ResNet

ResNet: ILSVRC 2015 winner

Revolution of Depth

AlexNet, 8 layers
(ILSVRC 2012)



VGG, 19 layers
(ILSVRC 2014)




ResNet, **152 layers**
(ILSVRC 2015)




ResNet

- The residual module
 - Introduce **skip or shortcut connections** (existing before in various forms in literature)
 - Make it easy for network layers to represent the identity mapping
 - For some reason, **need to skip at least two layers**



ResNet


Deeper residual module (bottleneck)




- Directly performing 3×3 convolutions with 256 feature maps at input and output:
 $256 \times 256 \times 3 \times 3 \sim 600K$ operations
- Using 1×1 convolutions to reduce 256 to 64 feature maps, followed by 3×3 convolutions, followed by 1×1 convolutions to expand back to 256 maps:
 $256 \times 64 \times 1 \times 1 \sim 16K$
 $64 \times 64 \times 3 \times 3 \sim 36K$
 $64 \times 256 \times 1 \times 1 \sim 16K$
Total: $\sim 70K$

ResNet


Reference paper



Batch Norm after add



No ReLU




ResNet

- Architectures for ImageNet problem:

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112			7×7, 64, stride 2		
conv2_x	56×56			3×3 max pool, stride 2		
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9


ResNet

- The architecture of the plain and residual networks were identical except for the skip connections
 - **Result: Going deeper makes things better!**





Why do ResNets work?

- ResNets seem to work because they **facilitate the training of deeper networks**
- Are **surprisingly robust** to layers being dropped or reordered
- Implicitly **ensembling shallower networks**
- Able to **learn unrolled iterative refinements**




“Understanding” Inception Module

Inception Module




Inception Module

- **Parallel paths with different receptive field sizes** and operations are meant to **capture sparse patterns of correlations** in the stack of feature maps




Inception Module


- Parallel paths **with different receptive field sizes** and operations are meant to capture sparse patterns of correlations **in the stack of feature maps**
- Use 1x1 convolutions for dimensionality **reduction before expensive convolutions**



GoogLeNet



GoogLeNet



Auxiliary classifier


GoogLeNet

- An alternative view:

type	patch size/ stride	output size	depth	#1×1	#3×3 reduce	#3×3	#5×5 reduce	#5×5	pool proj	params	ops
convolution	$7 \times 7 / 2$	$112 \times 112 \times 64$	1							2.7K	34M
max pool	$3 \times 3 / 2$	$56 \times 56 \times 64$	0								
convolution	$3 \times 3 / 1$	$56 \times 56 \times 192$	2		64	192				112K	360M
max pool	$3 \times 3 / 2$	$28 \times 28 \times 192$	0								
inception (3a)		$28 \times 28 \times 256$	2	64	96	128	16	32	32	159K	128M
inception (3b)		$28 \times 28 \times 480$	2	128	128	192	32	96	64	380K	304M
max pool	$3 \times 3 / 2$	$14 \times 14 \times 480$	0								
inception (4a)		$14 \times 14 \times 512$	2	192	96	208	16	48	64	364K	73M
inception (4b)		$14 \times 14 \times 512$	2	160	112	224	24	64	64	437K	88M
inception (4c)		$14 \times 14 \times 512$	2	128	128	256	24	64	64	463K	100M
inception (4d)		$14 \times 14 \times 528$	2	112	144	288	32	64	64	580K	119M
inception (4e)		$14 \times 14 \times 832$	2	256	160	320	32	128	128	840K	170M
max pool	$3 \times 3 / 2$	$7 \times 7 \times 832$	0								
inception (5a)		$7 \times 7 \times 832$	2	256	160	320	32	128	128	1072K	54M
inception (5b)		$7 \times 7 \times 1024$	2	384	192	384	48	128	128	1388K	71M
avg pool	$7 \times 7 / 1$	$1 \times 1 \times 1024$	0								
dropout (40%)		$1 \times 1 \times 1024$	0								
linear		$1 \times 1 \times 1000$	1							1000K	1M
softmax		$1 \times 1 \times 1000$	0								


Inception v2, v3

- Regularize training with **batch normalization**, reducing **importance of auxiliary classifiers**
- More variants of inception modules with **aggressive factorization of filters**



Inception v2, v3

- Increase the number of feature maps while **decreasing spatial resolution (pooling)**




What's new?

- **Batch Normalization (BN)** is used
- 1x1 convolution for dimensionality (z-axis) reduction
- **Average pooling introduced** in Inception module
- **Instead of 5x5 filters uses dual 3x3 filters**


Advantages

- **Reduce 90% computational parameters** compared to the AlexNet
- **Multiple receptive field for better stack(features)** of feature representation.
- **Achieve excellent performance** using limited number of parameters.

Inception-ResNet




Inception connections as introduced in C. Szegedy et al




Residual connections as introduced in He et al


Inception-Residual module



Inception v4



We're focusing on ImageNet




Summary: ILSVRC 2012-2015

Team	Year	Place	Error (top-5)	External data
SuperVision – Toronto (AlexNet, 7 layers)	2012	-	16.4%	no
SuperVision	2012	1st	15.3%	ImageNet 22k
Clarifai – NYU (7 layers)	2013	-	11.7%	no
Clarifai	2013	1st	11.2%	ImageNet 22k
VGG – Oxford (16 layers)	2014	2nd	7.32%	no
GoogLeNet (19 layers)	2014	1st	6.67%	no
ResNet (152 layers)	2015	1st	3.57%	
Human expert*			5.1%	

<http://karpathy.github.io/2014/09/02/what-i-learned-from-competing-against-a-convnet-on-imagenet/>

Accuracy vs. efficiency



Summary

- Introduce different CNN architectures, limitations and advantages
- General design principles of CNN models
- What's next?
 - Alternative of FCL, transfer Learning, and ensemble
 - Modern architectures:
 - Hybrid Networks
 - DenseNet
 - FactalNet

References

- <https://culurciello.github.io/tech/2016/06/04/nets.html>
- Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, [Gradient-based learning applied to document recognition](#), Proc. IEEE 86(11): 2278–2324, 1998.
- A. Krizhevsky, I. Sutskever, and G. Hinton, [ImageNet Classification with Deep Convolutional Neural Networks](#), NIPS 2012
- M. Zeiler and R. Fergus, [Visualizing and Understanding Convolutional Networks](#), ECCV 2014
- K. Simonyan and A. Zisserman, [Very Deep Convolutional Networks for Large-Scale Image Recognition](#), ICLR 2015
- M. Lin, Q. Chen, and S. Yan, [Network in network](#), ICLR 2014
- C. Szegedy et al., [Going deeper with convolutions](#), CVPR 2015
- C. Szegedy et al., [Rethinking the inception architecture for computer vision](#), CVPR 2016
- K. He, X. Zhang, S. Ren, and J. Sun, [Deep Residual Learning for Image Recognition](#), CVPR 2016